

Intro to Cryptography

CS 161 Spring 2024 - Lecture 5

Last Time: Memory Safety Mitigations

Computer Science 161

- Memory-safe languages
 - Using a memory-safe language (e.g. Python, Java) stops all memory safety vulnerabilities.
 - Why use a non-memory-safe language?
 - Commonly-cited reason, but mostly a myth: Performance
 - Real reason: Legacy, existing code
- Writing memory-safe code
 - Carefully write and reason about your code to ensure memory safety in a non-memory-safe language
 - Requires programmer discipline, and can be tedious sometimes
- Building secure software
 - Use tools for analyzing and patching insecure code
 - Test your code for memory safety vulnerabilities
 - Keep any external libraries updated for security patches

Last Time: Memory Safety Mitigations

Computer Science 161

- Mitigation: **Non-executable pages**

- Make portions of memory either executable or writable, but not both
- Defeats attacker writing shellcode to memory and executing it
- Subversions
 - **Return-to-libc**: Execute an existing function in the C library
 - **Return-oriented programming (ROP)**: Create your own code by chaining together small gadgets in existing library code

- Mitigation: **Stack canaries**

- Add a sacrificial value on the stack. If the canary has been changed, someone's probably attacking our system
- Defeats attacker overwriting the RIP with address of shellcode
- Subversions
 - An attacker can write around the canary
 - The canary can be leaked by another vulnerability (e.g. format string vulnerability)
 - The canary can be brute-forced by the attacker

Last Time: Memory Safety Mitigations

Computer Science 161

- Mitigation: **Pointer authentication**
 - When storing a pointer in memory, replace the unused bits with a pointer authentication code (PAC). Before using the pointer in memory, check if the PAC is still valid
 - Defeats attacker overwriting the RIP (or any pointer) with address of shellcode
- Mitigation: **Address space layout randomization (ASLR)**
 - Put each segment of memory in a different location each time the program is run
 - Defeats attacker knowing the address of shellcode
 - Subversions
 - Leak addresses with another vulnerability
 - Brute-force attack to guess the addresses
- Combining mitigations
 - Using multiple mitigations usually forces the attacker to find multiple vulnerabilities to exploit the program (defense-in-depth)

What is cryptography?

Textbook Chapter 5.1

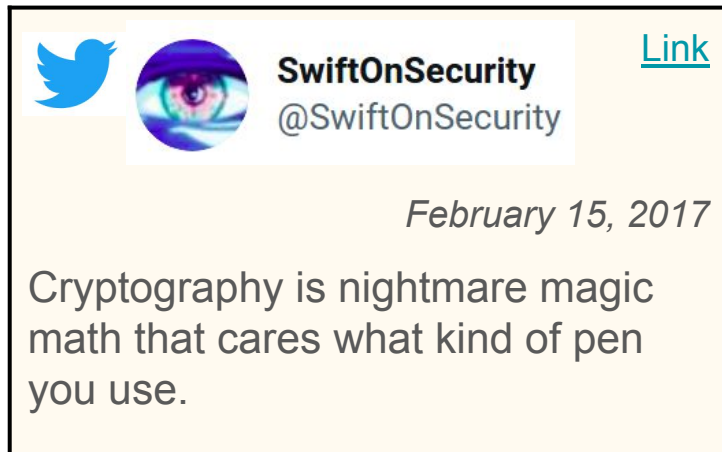
What is cryptography?

- Older definition: The study of secure communication over insecure channels
- Newer definition: Provide rigorous guarantees about the data and computation in the presence of an attacker
 - Not just *confidentiality* but also *integrity* and *authenticity* (we'll see these definitions today)
- Modern cryptography involves a lot of math
 - We'll review any necessary CS 70 prerequisites as they come up

Don't try this at home!

Computer Science 161

- We will teach you the basic building blocks of cryptography, but you should never try to write your own cryptographic algorithms
- It's very easy to make a mistake that makes your code insecure
 - Required formal rigor and tricky edge cases that we won't cover
 - One small bug could compromise the security of your code
- Instead, use existing well-vetted cryptographic libraries
 - This portion of the class is as much about making you a good *consumer* of cryptography



Don't try this at home!


- In summer 2020, CS 61A wrote a program to distribute online exams
- However, when writing cryptographic code, they used a secure algorithm in an insecure way
- Because of their mistake, it was possible to see exam questions before they were released!
 - Exam leakage was reported, but we never found out if anyone actually attacked the insecure scheme
- The TAs who wrote this code were former CS 161 students!
- **Takeaway:** Do not write your own crypto code!

Definitions

Textbook Chapter 5.3–5.9

Meet Alice, Bob, Eve, and Mallory

- Alice and Bob: The main characters trying to send messages to each other over an insecure communication channel
 - Carol and Dave can also join the party later
- Eve: An **eavesdropper** who can read any data sent over the channel
 - Also called an honest-but-curious attacker
 - Does not tamper with data or code
- Mallory: A **manipulator** who can read and modify any data sent over the channel
 - Also called a malicious attacker



Meet Alice, Bob, Eve, and Mallory

- We often describe cryptographic problems using a common cast of characters
- One scenario:
 - Alice wants to send a message to Bob.
 - However, Eve is going to *eavesdrop* on the communication channel.
 - How does Alice send the message to Bob without Eve learning about the message?
- Another scenario:
 - Bob wants to send a message to Alice.
 - However, Mallory is going to *tamper* with the communication channel.
 - How does Bob send the message to Alice without Mallory changing the message?

Three Main Goals of Cryptography

- In cryptography, there are three common properties that we want on our data
- **Confidentiality**: An adversary cannot *read* our messages.
- **Integrity**: An adversary cannot *change* our messages without being detected.
- **Authenticity**: I can prove that this message came from the person who claims to have written it.
 - Integrity and authenticity are closely related properties...
 - Before I can prove that a message came from a certain person, I have to prove that the message wasn't changed!
 - ... but they're not identical properties
 - Later we'll see some edge cases

Keys

- The most basic building block of any cryptographic scheme: The **key**
- Properly chosen and guarded keys “power” the security of our cryptographic algorithms
- Two models of keys:
 - **Symmetric key model:** Alice and Bob both know the value of the same secret key.
 - **Asymmetric key model:** A user has two keys, a secret key and a public key.
 - Example: You might remember RSA encryption from CS 70

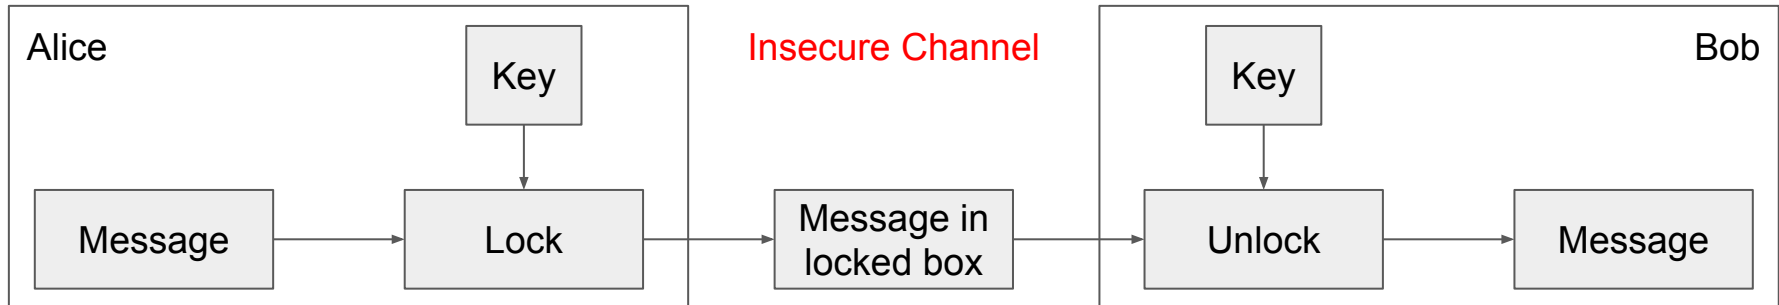


Security Principle: Kerckhoff's Principle

- This principle is closely related to Shannon's Maxim
 - Don't use security through obscurity. Assume the attacker knows the system.
- Kerckhoff's principle says:
 - Cryptosystems should remain secure even when the attacker knows all internal details of the system
 - The key should be the only thing that must be kept secret
 - The system should be designed to make it easy to change keys that are leaked (or suspected to be leaked)
 - If your secrets are leaked, it is usually a lot easier to change the key than to replace every instance of the running software
- Our assumption: The attacker knows all the algorithms we use. The only information the attacker is missing is the secret key(s).

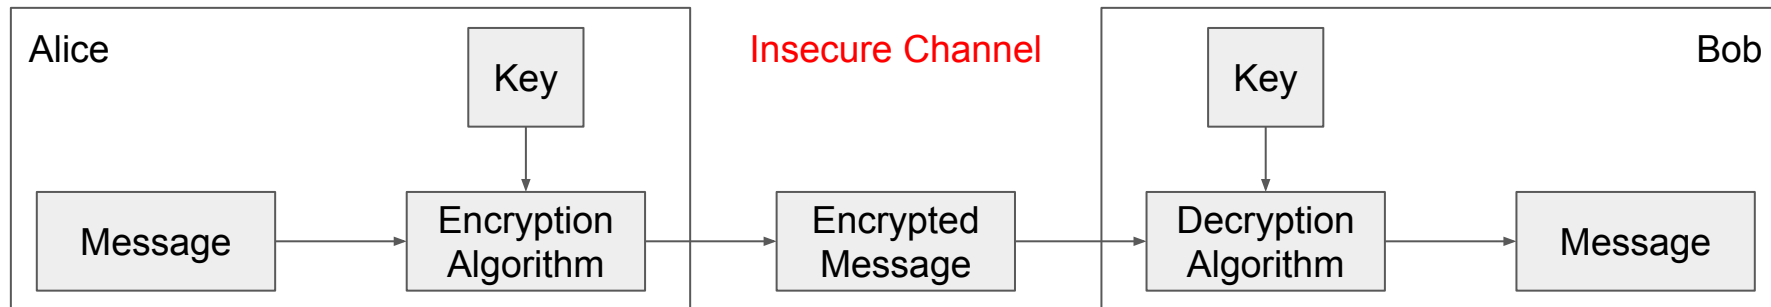
Confidentiality

- **Confidentiality:** An adversary cannot *read* our messages.
- **Analogy:** Locking and unlocking the message
 - Alice uses the key to lock the message in a box
 - Alice sends the message (locked in the box) over the insecure channel
 - Eve sees the locked box, but cannot access the message without the key
 - Bob receives the message (locked in the box) and uses the key to unlock the message




Confidentiality

- Schemes provide confidentiality by **encrypting** messages
 - Alice uses the key to **encrypt** the message: Change the message into a scrambled form
 - Alice sends the encrypted message over the insecure channel
 - Eve sees the encrypted message, but cannot figure out the original message without the key
 - Bob receives the encrypted message and uses the key to **decrypt** the message back into its original form



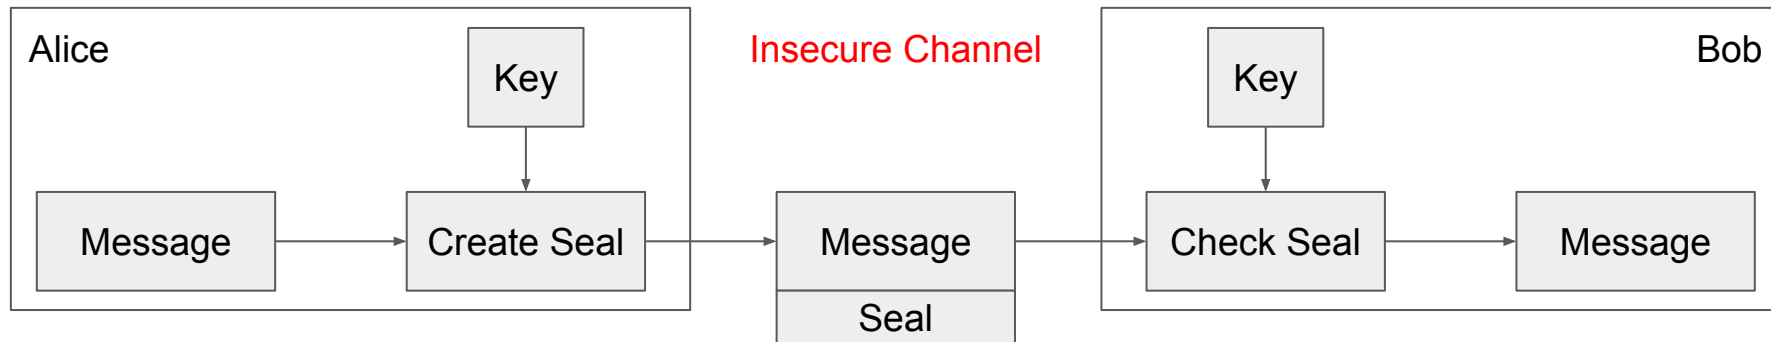
Confidentiality

- **Plaintext:** The original message
- **Ciphertext:** The encrypted message




Integrity (and Authenticity)

- **Integrity:** An adversary cannot *change* our messages without being detected.
- **Analogy:** Adding a seal on the message
 - Alice uses the key to add a special seal on the message (e.g. puts tape on the envelope)
 - Alice sends the message and the seal over the insecure channel
 - If Mallory tampers with the message, she'll break the seal (e.g. break the tape on the envelope)
 - Without the key, Mallory cannot create her own seal
 - Bob receives the message and the seal and checks that the seal has not been broken



Integrity (and Authenticity)

- Schemes provide integrity by adding a **tag** or **signature** on messages
 - Alice uses the key to generate a special tag for the message
 - Alice sends the message and the tag over the insecure channel
 - If Mallory tampers with the message, the tag will no longer be valid
 - Bob receives the message and the tag and checks that the tag is still valid
- More on integrity in a future lecture



Threat Models

- What if Eve can do more than eavesdrop?
- Real-world schemes are often vulnerable to more sophisticated attackers, so cryptographers have created more sophisticated threat models too
- Some threat models for analyzing confidentiality:

	Can Eve trick Alice into encrypting messages of Eve's choosing?	Can Eve trick Bob into decrypting messages of Eve's choosing?
Ciphertext-only	No	No
Chosen-plaintext	Yes	No
Chosen-ciphertext	No	Yes
Chosen plaintext-ciphertext	Yes	Yes

Threat Models

- In this class, we'll explain the chosen plaintext attack model
- In practice, cryptographers use the chosen plaintext-ciphertext model
 - It's the most powerful
 - It can actually be defended against

	Can Eve trick Alice into encrypting messages of Eve's choosing?	Can Eve trick Bob into decrypting messages of Eve's choosing?
Ciphertext-only	No	No
Chosen-plaintext	Yes	No
Chosen-ciphertext	No	Yes
Chosen plaintext-ciphertext	Yes	Yes

Cryptography Roadmap

Computer Science 161

	Symmetric-key	Asymmetric-key
Confidentiality	<ul style="list-style-type: none">• One-time pads• Block ciphers with chaining modes (e.g. AES-CBC)• Stream ciphers	<ul style="list-style-type: none">• RSA encryption• ElGamal encryption
Integrity, Authentication	<ul style="list-style-type: none">• MACs (e.g. HMAC)	<ul style="list-style-type: none">• Digital signatures (e.g. RSA signatures)

- Hash functions
- Pseudorandom number generators
- Public key exchange (e.g. Diffie-Hellman)
- Key management (certificates)
- Password management

Symmetric-Key Encryption



Textbook Chapter 6.1

Cryptography Roadmap

Computer Science 161

	Symmetric-key	Asymmetric-key
Confidentiality	<ul style="list-style-type: none">• One-time pads• Block ciphers with chaining modes (e.g. AES-CBC)• Stream ciphers	<ul style="list-style-type: none">• RSA encryption• ElGamal encryption
Integrity, Authentication	<ul style="list-style-type: none">• MACs (e.g. HMAC)	<ul style="list-style-type: none">• Digital signatures (e.g. RSA signatures)

- Hash functions
- Pseudorandom number generators
- Public key exchange (e.g. Diffie-Hellman)

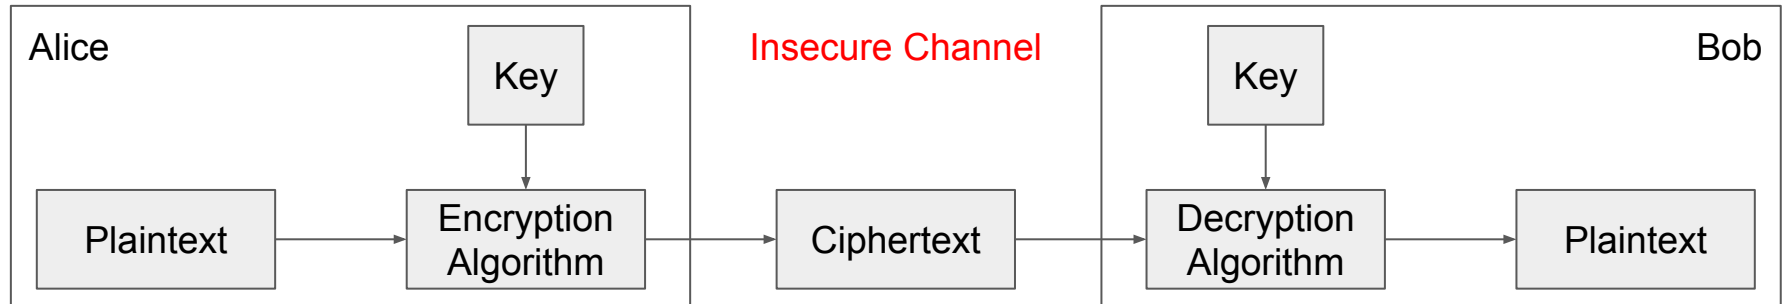
- Key management (certificates)
- Password management

Symmetric-Key Encryption

- The next few schemes are symmetric-key encryption schemes
 - **Encryption schemes** aim to provide *confidentiality* (but not integrity or authentication)
 - **Symmetric-key** means Alice and Bob share the same secret key that the attacker doesn't know
 - Don't worry about how Alice and Bob share the key for now
- For modern schemes, we're going to assume that messages are *bitstrings*
 - **Bitstring**: A sequence of bits (0 or 1), e.g. 11010101001001010
 - Text, images, etc. can be converted into bitstrings before encryption, so bitstrings are a useful abstraction. After all, everything in a computer is just a sequence of bits!

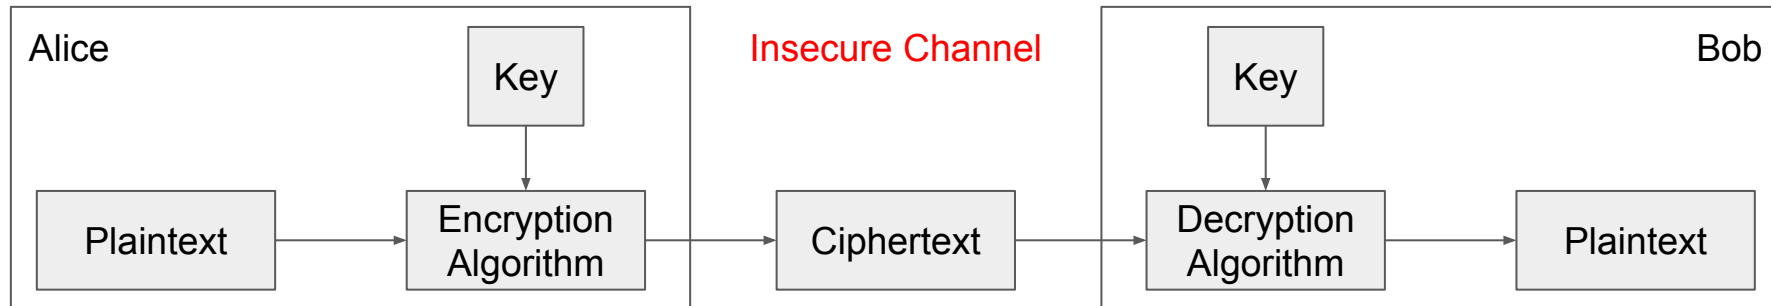
Symmetric-Key Encryption: Definition

- A symmetric-key encryption scheme has three algorithms:
 - $\text{KeyGen}() \rightarrow K$: Generate a key K
 - $\text{Enc}(K, M) \rightarrow C$: Encrypt a **plaintext** M using the key K to produce **ciphertext** C
 - $\text{Dec}(K, C) \rightarrow M$: Decrypt a ciphertext C using the key K



Symmetric-Key Encryption: Definition

- What properties do we want from a symmetric encryption scheme?
 - **Correctness:** Decrypting a ciphertext should result in the message that was originally encrypted
 - $\text{Dec}(K, \text{Enc}(K, M)) = M$ for all $K \leftarrow \text{KeyGen}()$ and M
 - **Efficiency:** Encryption/decryption algorithms should be fast: >1 Gbps on a standard computer
 - **Security:** Confidentiality



Defining Confidentiality

- Recall our definition of confidentiality from earlier: “An adversary cannot read our messages”
 - This definition isn’t very specific
 - What if Eve can read the first half of Alice’s message, but not the second half?
 - What if Eve figures out that Alice’s message starts with “Dear Bob”?
 - This definition doesn’t account for prior knowledge
 - What if Eve already knew that Alice’s message ends in “Sincerely, Alice”?
 - What if Eve knows that Alice’s message is “BUY!” or “SELL” but doesn’t know which?

Q: How would you define confidentiality?




Defining Confidentiality

- A better definition of confidentiality: The ciphertext should not give the attacker *any additional information* about the plaintext.
- Let's design an experiment/security game to test our definition

Security game: first attempt at confidentiality


1. Eve issues a pair of plaintexts M_0 and M_1 to Alice of the same length
2. Alice randomly chooses either M_0 or M_1 to encrypt and sends the encryption back
 - a. Alice does not tell Eve which one was encrypted!
3. Eventually, Eve outputs a guess as to whether Alice encrypted M_0 or M_1

Q: If the scheme provides confidentiality, what chance does the attacker have to guess b ?



Security game: intuition

- If the scheme is secure Eve can only guess with probability $1/2$, which is no different than if Eve hadn't sent the ciphertext at all
- In other words: the ciphertext gave Eve no *additional* information about which plaintext was sent!

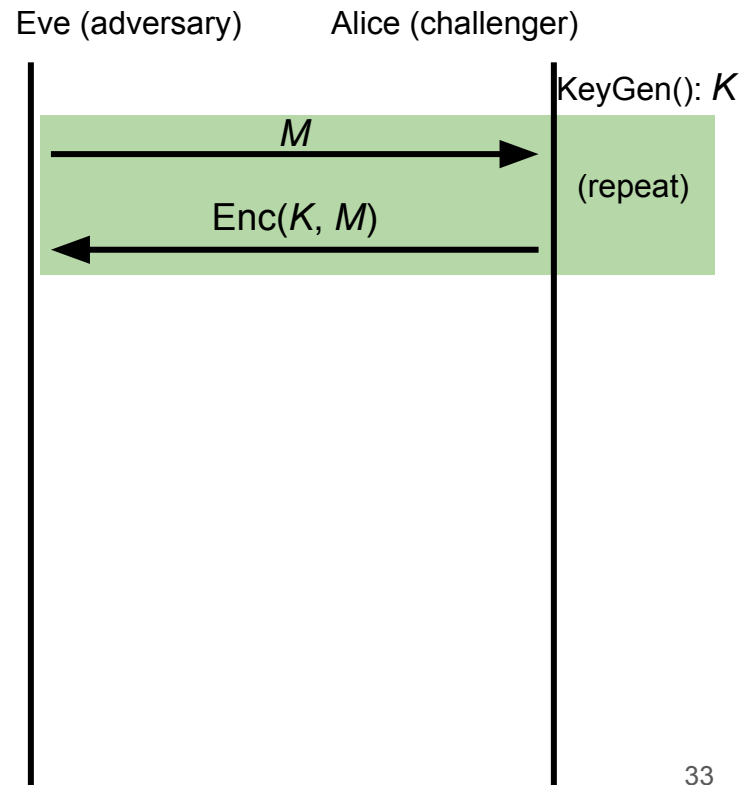


Defining Confidentiality: IND-CPA

- Recall our threat model: Eve can also perform a **chosen plaintext attack**
 - Eve can trick Alice into encrypting arbitrary messages of Eve's choice
 - We can adapt our experiment to account for this threat model
- A better definition of confidentiality: Even if Eve is able to trick Alice into encrypting messages, Eve can still only guess what message Alice sent with probability $1/2$.
 - This definition is called **IND-CPA** (indistinguishability under chosen plaintext attack)
- Cryptographic properties are often defined in terms of “games” that an adversary can either “win” or “lose”
 - We will use one to define confidentiality precisely

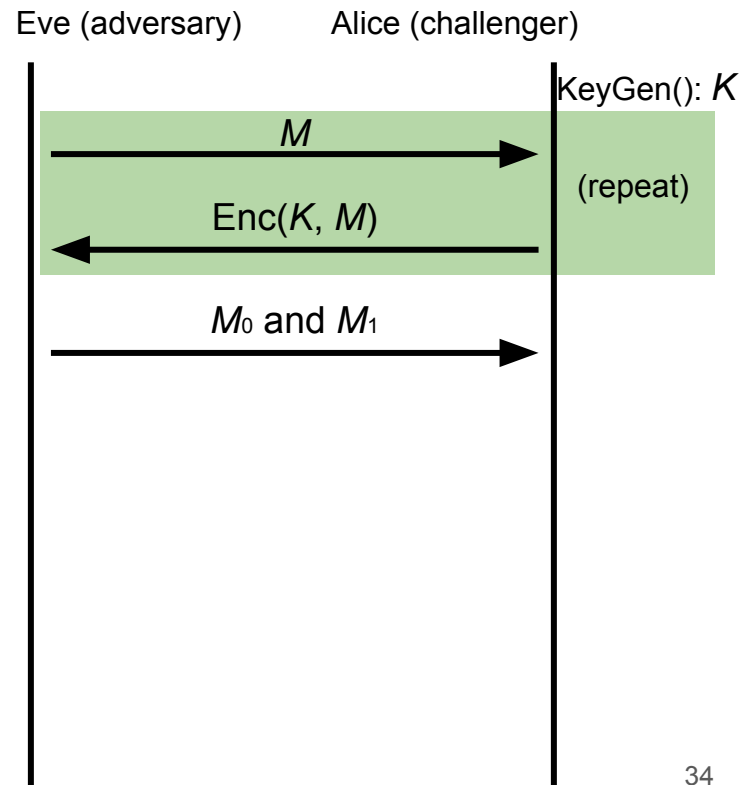
Defining Confidentiality: IND-CPA

1. Eve may choose plaintexts to send to Alice and receives their ciphertexts



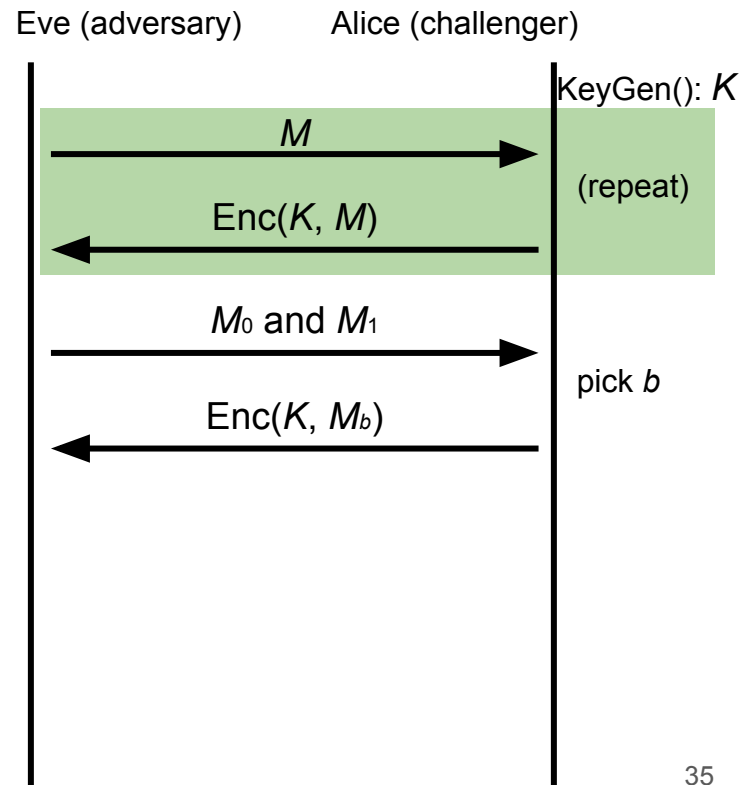
Defining Confidentiality: IND-CPA

1. Eve may choose plaintexts to send to Alice and receives their ciphertexts
2. Eve issues a pair of plaintexts M_0 and M_1 to Alice



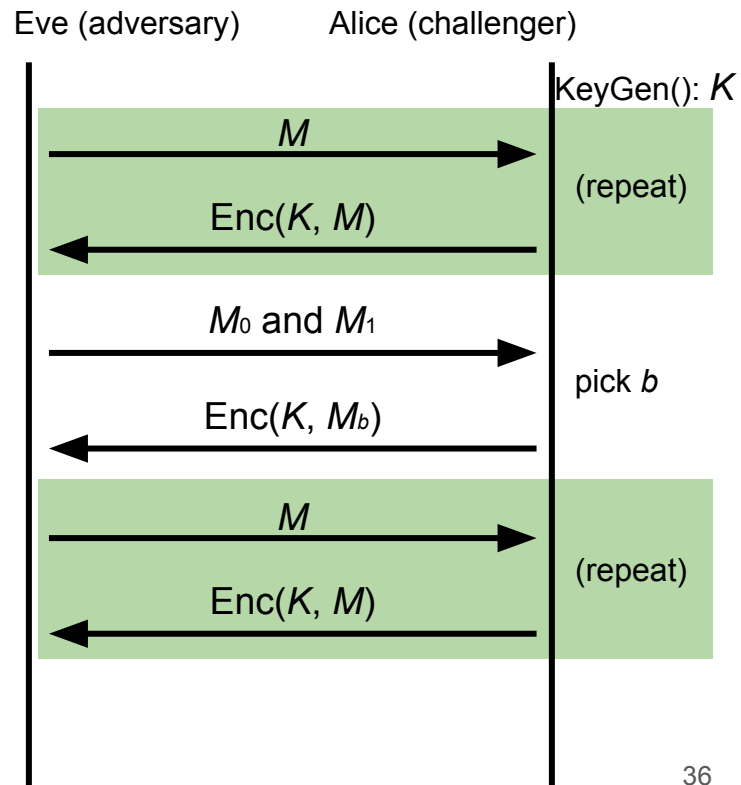
Defining Confidentiality: IND-CPA

1. Eve may choose plaintexts to send to Alice and receives their ciphertexts
2. Eve issues a pair of plaintexts M_0 and M_1 to Alice
3. Alice randomly chooses either M_0 or M_1 to encrypt and sends the encryption back
 - a. Alice does not tell Eve which one was encrypted!



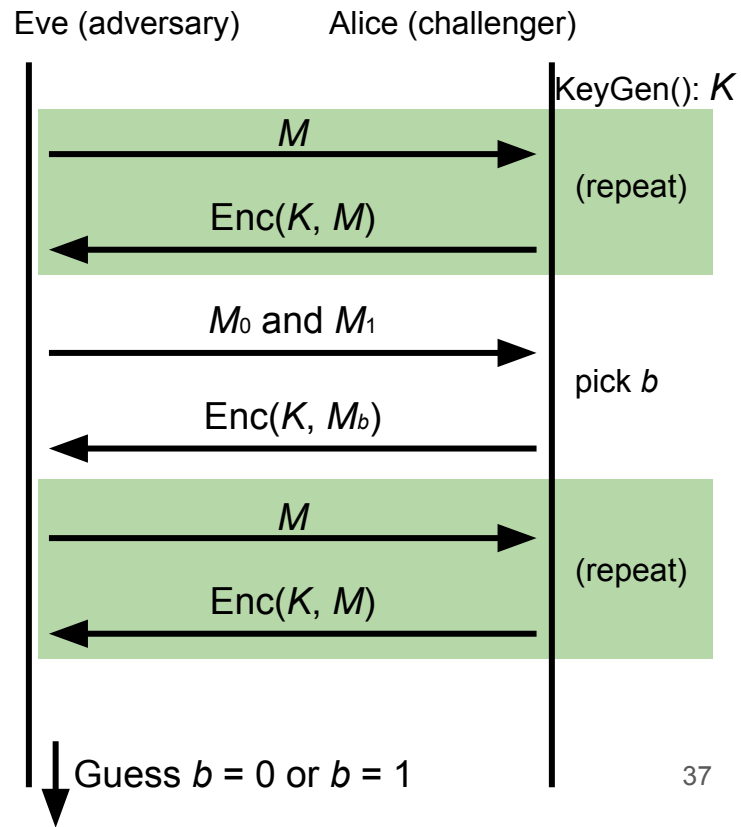
Defining Confidentiality: IND-CPA

1. Eve may choose plaintexts to send to Alice and receives their ciphertexts
2. Eve issues a pair of plaintexts M_0 and M_1 to Alice
3. Alice randomly chooses either M_0 or M_1 to encrypt and sends the encryption back
 - a. Alice does not tell Eve which one was encrypted!
4. Eve may again choose plaintexts to send to Alice and receives their ciphertexts



Defining Confidentiality: IND-CPA

1. Eve may choose plaintexts to send to Alice and receives their ciphertexts
2. Eve issues a pair of plaintexts M_0 and M_1 to Alice
3. Alice randomly chooses either M_0 or M_1 to encrypt and sends the encryption back
 - a. Alice does not tell Eve which one was encrypted!
4. Eve may again choose plaintexts to send to Alice and receives their ciphertexts
5. Eventually, Eve outputs a guess as to whether Alice encrypted M_0 or M_1



Defining Confidentiality: IND-CPA

- If Eve correctly guesses which message Alice encrypted, then Eve wins. Otherwise, she loses.
- How does Eve guess whether M_0 or M_1 was encrypted? What strategy does she use?
 - We don't *assume* she uses a particular strategy; Eve represents all possible strategies
- Proving insecurity: There exists at least *one* strategy that can win the IND-CPA game with probability $> 1/2$
 - $1/2$ is the probability of winning by random guessing
 - If you can be better than random, then the ciphertext has leaked information, and Eve is able to learn it and use it to gain an advantage!
- Proving security: For *all* (polynomial-time) attackers/Eve-s, the probability of winning the IND-CPA game is at most $1/2 + \text{negl}$

Edge Cases: Length

- Cryptographic schemes are (usually) allowed to leak the length of the message
 - To hide length: All messages must always be the same length
 - 16-byte messages: We can't encrypt large messages (images, videos, etc.)!
 - 1-GB messages: Sending small messages (text, Tweets, etc.) needs 1 GB of bandwidth!
 - This is unpractical
 - Applications can choose to hide length by *padding* their own messages to the maximum possible length before encrypting
- In the IND-CPA game: M_0 and M_1 must be the same length
 - To break IND-CPA, Eve must learn something other than message length



Edge Cases: Attacker Runtime

- Some schemes are theoretically vulnerable, but secure in any real-world setting
 - If an attack takes longer than the life of the solar system to complete, it probably won't happen!
 - Or if it would require a computer made out of a literal galaxy worth of science-fiction nanotech
- In the IND-CPA game: Eve is limited to a practical runtime
 - One common practical limit: Eve is limited to polynomial runtime algorithms (no exponential-time algorithms)



Edge Cases: Negligible Advantage

- Sometimes it's possible for Eve to win with probability $1/2 + 1/2^{128}$
 - This probability is greater than $1/2$, but it's so close to $1/2$ that it's as good as $1/2$.
 - Eve's advantage is so small that she can't use it for any practical attacks
 - 2^{128} is larger than the total number of atoms in the universe
- In the IND-CPA game: The scheme is secure even if Eve can win with probability $\leq 1/2 + \epsilon$, where ϵ is *negligible*
 - The actual mathematical definition of negligible is out of scope
 - Example: $1/2 + 1/2^{128}$: Negligible advantage
 - Example: $2/3$: Non-negligible advantage



Edge Cases: Negligible Advantage

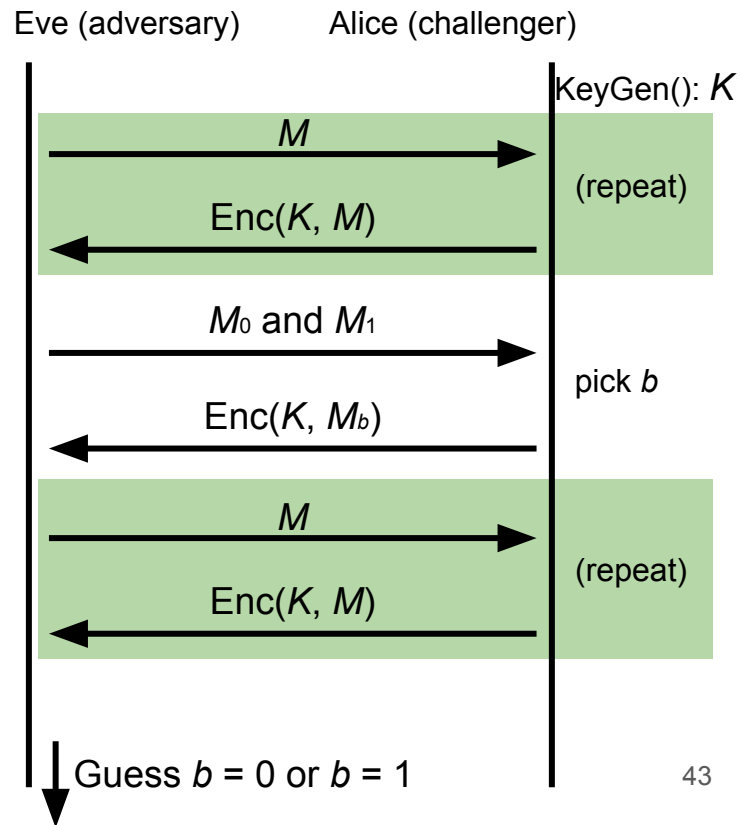
- Defining negligibility mathematically:
 - Advantage of the adversary should be exponentially small, based on the security parameters of the algorithm
 - Example: For an encryption scheme with a k -bit key, the advantage should be $O(1/2^k)$
- Defining negligibility practically:
 - A $1/2^{128}$ probability is currently unlikely
 - A $1/2^{20}$ probability is fairly likely
 - “One in a million events happen every day in New York City”
 - In between these extremes, it can be messy
 - Different algorithms run faster or slower and have their own security parameters
 - Computers get more powerful over time
 - Recall: Know your threat model!
- **Takeaway:** For now, 2^{80} is a reasonable threshold, but this will change over time!

IND-CPA: Putting it together

1. Eve may choose plaintexts to send to Alice and receives their ciphertexts
2. Eve issues a pair of plaintexts M_0 and M_1 to Alice
3. Alice randomly chooses either M_0 or M_1 to encrypt and sends the encryption back
 - Alice does not tell Eve which one was encrypted!
4. Eve may again choose plaintexts to send to Alice and receives their ciphertexts
5. Eventually, Eve outputs a guess as to whether Alice encrypted M_0 or M_1



- An encryption scheme is IND-CPA secure if for all polynomial time attackers Eve:
 - Eve can win with probability $\leq 1/2 + \epsilon$, where ϵ is *negligible*.



A Brief History of Cryptography



Textbook Chapter 5.2

Cryptography by Hand: Caesar Cipher

Computer Science 161

- One of the earliest cryptographic schemes was the **Caesar cipher**
 - Used by Julius Caesar over 2,000 years ago
- **KeyGen():**
 - Choose a key K randomly between 0 and 25
- **Enc(K, M):**
 - Replace each letter in M with the letter K positions later in the alphabet
 - If $K = 3$, plaintext DOG becomes GRJ
- **Dec(K, C):**
 - Replace each letter in C with the letter K positions earlier in the alphabet
 - If $K = 3$, ciphertext GRJ becomes DOG

$K = 3$			
M	C	M	C
A	D	N	Q
B	E	O	R
C	F	P	S
D	G	Q	T
E	H	R	U
F	I	S	V
G	J	T	W
H	K	U	X
I	L	V	Y
J	M	W	Z
K	N	X	A
L	O	Y	B
M	P	Z	C

Cryptography by Hand: Attacks on the Caesar Cipher

Computer Science 161

- Eve sees the ciphertext JCKN ECGUCT, but doesn't know the key K
- If you were Eve, how would you try to break this algorithm?
- Brute-force attack: Try all 26 possible keys!
- Use existing knowledge: Assume that the message is in English

+1 IBJM DBFTBS

+2 HAIL CAESAR

+3 GZHK BZDRZQ

+4 FYGJ AYCQYP

+5 EXFI ZXBPXO

+6 DWEH YWAOWN

+7 CVDG XVZNVN

+8 BUCF WUYMUL

+9 ATBE VTXLTK

+10 ZSAD USWKSJ

+11 YRZC TRVJRI

+12 XQYB SQUIQH

+13 WPXA RPTHPI

+14 VOWZ QOSGOF

+15 UNVY PNRFNE

+16 TMUX OMQEMD

+17 SLTW NLPDLC

+18 RKSV MKOCKB

+19 QJRU LJNBJA

+20 PIQT KIMAIZ

+21 OHPS JHLZHY

+22 NGOR IGKYGX

+23 MFNQ HFJXFW

+24 LEMP GEIWEV

+25 KDLO FDHVDU

Cryptography by Hand: Attacks on the Caesar Cipher

Computer Science 161

- Eve sees the ciphertext JCKN ECGUCT, but doesn't know the key K
- Chosen-plaintext attack: Eve tricks Alice into encrypting plaintext of her choice
 - Eve sends a message $M = \text{AAA}$ and receives $C = \text{CCC}$
 - Eve can deduce the key: C is 2 letters after A, so $K = 2$
 - Eve has the key, so she can decrypt the ciphertext

Cryptography by Hand: Substitution Cipher

Computer Science 161

- A better cipher: create a mapping of each character to another character.
 - Example: A = N, B = Q, C = L, D = Z, etc.
 - Unlike the Caesar cipher, the shift is no longer constant!
- **KeyGen():**
 - Generate a random, one-to-one mapping of characters
- **Enc(K , M):**
 - Map each letter in M to the output according to the mapping K
- **Dec(K , C):**
 - Map each letter in C to the output according to the *reverse* of the mapping K

K			
M	C	M	C
A	N	N	G
B	Q	O	P
C	L	P	T
D	Z	Q	A
E	K	R	J
F	R	S	O
G	V	T	D
H	U	U	I
I	E	V	C
J	S	W	F
K	B	X	M
L	W	Y	X
M	Y	Z	H

Cryptography by Hand: Attacks on Substitution Ciphers

Computer Science 161

- Does the brute-force attack still work?
 - There are $26! \approx 2^{88}$ possible mappings to try
 - Too much for most modern computers... for now
- How about the chosen-plaintext attack?
 - Trick Alice into encrypting ABCDEFGHIJKLMNOPQRSTUVWXYZ, and you'll get the whole mapping!
- Another strategy: *cryptanalysis*
 - The most common english letters in text are E, T, A, O, I, N

K			
M	C	M	C
A	N	N	G
B	Q	O	P
C	L	P	T
D	Z	Q	A
E	K	R	J
F	R	S	O
G	V	T	D
H	U	U	I
I	E	V	C
J	S	W	F
K	B	X	M
L	W	Y	X
M	Y	Z	H

Takeaways

Computer Science 161



- Cryptography started with paper-and-pencil algorithms (Caesar cipher)
- Then cryptography moved to machines (Enigma)
- Finally, cryptography moved to computers (which we're about to study)
- Hopefully you gained some intuition for some of the cryptographic definitions



Cryptography by Machines: Enigma

Computer Science 161

- A mechanical encryption machine used by the Germans in WWII




Enigma Operating Principle: Rotor Machine

- The encryption core was composed of 3 or 4 rotors
 - Each rotor was a fixed permutation (e.g. A maps to F, B maps to Q...)
 - And the end was a "reflector", a rotor that sent things backwards
 - Plus a fixed-permutation plugboard
- A series of rotors were arranged in a sequence
 - Each keypress would generate a current from the input to one light for the output
 - Each keypress also advanced the first rotor
 - When the first rotor makes a full rotation, the second rotor advances one step
 - When the second rotor makes a full rotation, the third rotor advances once step

Cryptography by Machines: Enigma

Computer Science 161

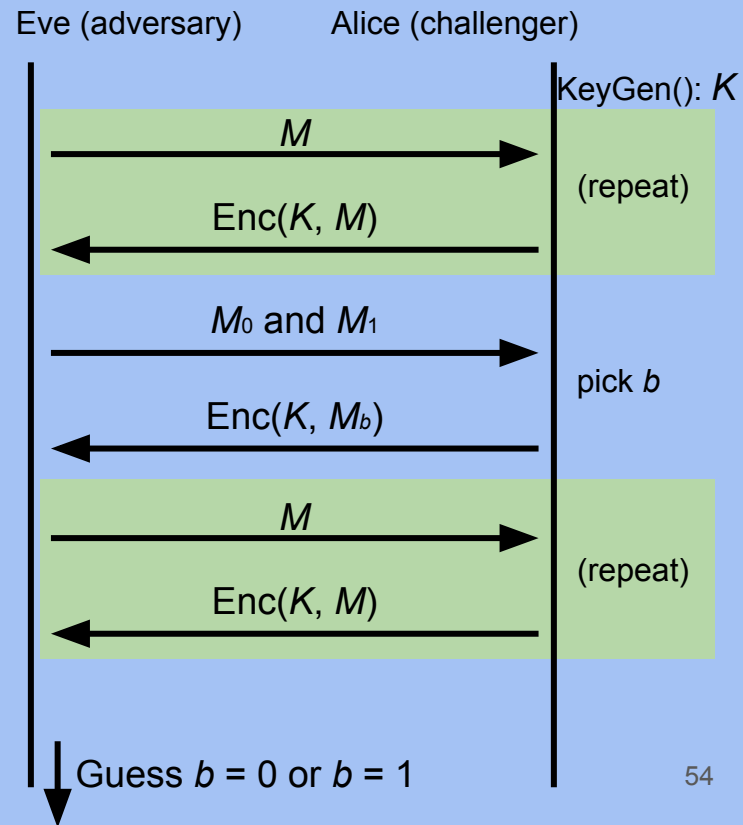
- **KeyGen():**
 - Choose rotors, rotor orders, rotor positions, and plugboard settings
 - 158,962,555,217,826,360,000 possible keys
- **Enc(K , M) and Dec(K , C):**
 - Input the rotor settings K into the Enigma machine
 - Press each letter in the input, and the lampboard will light up the corresponding output letter
 - Encryption and decryption are the same algorithm!
- Germans believed that Enigma was an “unbreakable code”



Cryptography by Machines: Enigma

Computer Science 161

- Enigma has a significant weakness: a letter never maps to itself!
 - No rotor maps a letter to itself
 - The reflector never maps a letter to itself
 - This property is necessary for Enigma's mechanical system to work
- What pair of messages should Eve send to Alice in the challenge phase?
 - Send $M_0 = A^k$, $M_1 = B^k$
 - M_0 is a string of k 'A' characters, M_1 is a string of k 'B' characters
- How can Eve probably know which message Alice encrypted?



Cryptography by Machines: Attack on Enigma

Computer Science 161

- Polish and British cryptographers built BOMBE, a machine to brute-force Enigma keys
- Why was Enigma breakable?
 - Kerckhoff's principle: The Allies stole Enigma machines, so they knew the algorithm
 - Known plaintext attacks: the Germans often sent predictable messages (e.g. the weather report every morning)
 - Chosen plaintext attacks: the Allies could trick the Germans into sending a message (e.g. "newly deployed minefield")
 - Brute-force: BOMBE would try many keys until the correct one was found
 - Plus a weakness: You'd be able to try multiple keys with the same hardware configuration




BOMBE machine

Cryptography by Machines: Legacy of Enigma

Computer Science 161

- Alan Turing, one of the cryptographers who broke Enigma, would go on to become one of the founding fathers of computer science
- Most experts agree that the Allies breaking Enigma shortened the war in Europe by about a year



Alan Turing

Cryptography by Computers

Computer Science 161

- The modern era of cryptography started after WWII, with the work of Claude Shannon
- “New Directions in Cryptography” (1976) showed how number theory can be used in cryptography
 - Its authors, Whitfield Diffie and Martin Hellman, won the Turing Award in 2015 for this paper
- This is the era of cryptography we’ll be focusing on



One of these is Diffie, and
the other one is Hellman.

