

Block Cipher Chaining Modes (cont'd) & Cryptographic Hashes

CS 161 Spring 2023 - Lecture 7

Announcements

Computer Science 161

- Project 1 (Q1–Q7, plus write-up) is due Friday, February 9th at 11:59 PM PT.
 - When you are writing your write-up please explain the vulnerability and your exploit in a detailed manner. Your write-up should include the logical steps you take when you are coming up with your exploit as well as the magic numbers you found using GDB.
 - Please make sure you fill out the OH template when you are making your ticket. The tickets that don't provide this template will not be taken.
- Homework 2 is due Friday, February 9th at 11:59 PM PT.
- The midterm is on Thursday, February 29th from 7:00-9:00 PM PT.
 - If you would like to request an alternate exam time or remote exam, or have DSP accommodations or any special requests, please fill out [the Exam Logistics Form](#) by Monday, February 19, 11:59 PM PT.

Last Time: Block Ciphers

Computer Science 161

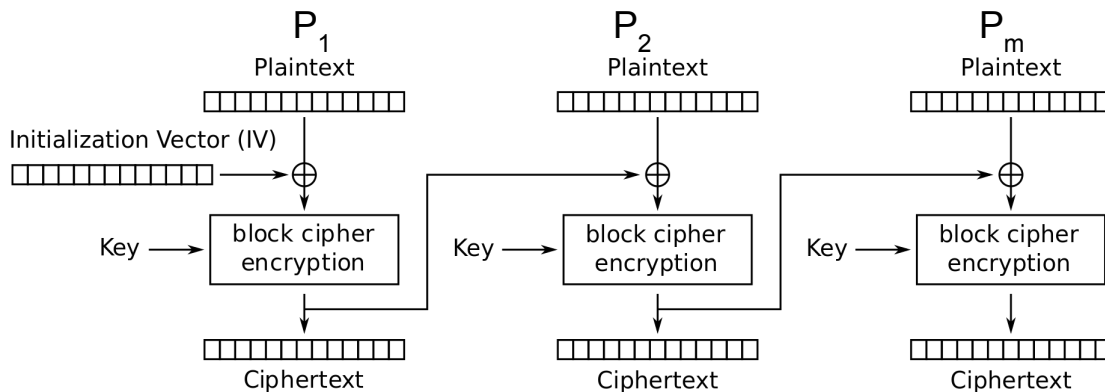
- Encryption: input a k -bit key and n -bit plaintext, receive n -bit ciphertext
- Decryption: input a k -bit key and n -bit ciphertext, receive n -bit plaintext
- Correctness: when the key is fixed, $E_k(M)$ should be bijective
- Security
 - Without the key, $E_k(m)$ is computationally indistinguishable from a random permutation
 - Brute-force attacks take astronomically long and are not possible
- Efficiency: algorithms use XORs and bit-shifting (very fast)
- Implementation: AES is the modern standard
- Issues
 - Not IND-CPA secure because they're deterministic
 - Can only encrypt n -bit messages

Block Cipher Modes of Operation: Summary

- ECB mode: Deterministic, so not IND-CPA secure
- CBC mode
 - IND-CPA secure, assuming no IV reuse
 - Encryption is not parallelizable
 - Decryption is parallelizable
 - Must pad plaintext to a multiple of the block size
 - IV reuse leads to leaking the existence of identical blocks at the start of the message

Recall: CBC Mode

- We've just designed **cipher block chaining (CBC) mode**
- $C_i = E_K(M_i \oplus C_{i-1})$; $C_0 = IV$
- $\text{Enc}(K, M)$:
 - Split M in m plaintext blocks $P_1 \dots P_m$ each of size n
 - Choose a random IV
 - Compute and output (IV, C_1, \dots, C_m) as the overall ciphertext
- How do we decrypt?




Cipher Block Chaining (CBC) mode encryption

CTR Mode Scratchpad: Let's design it together

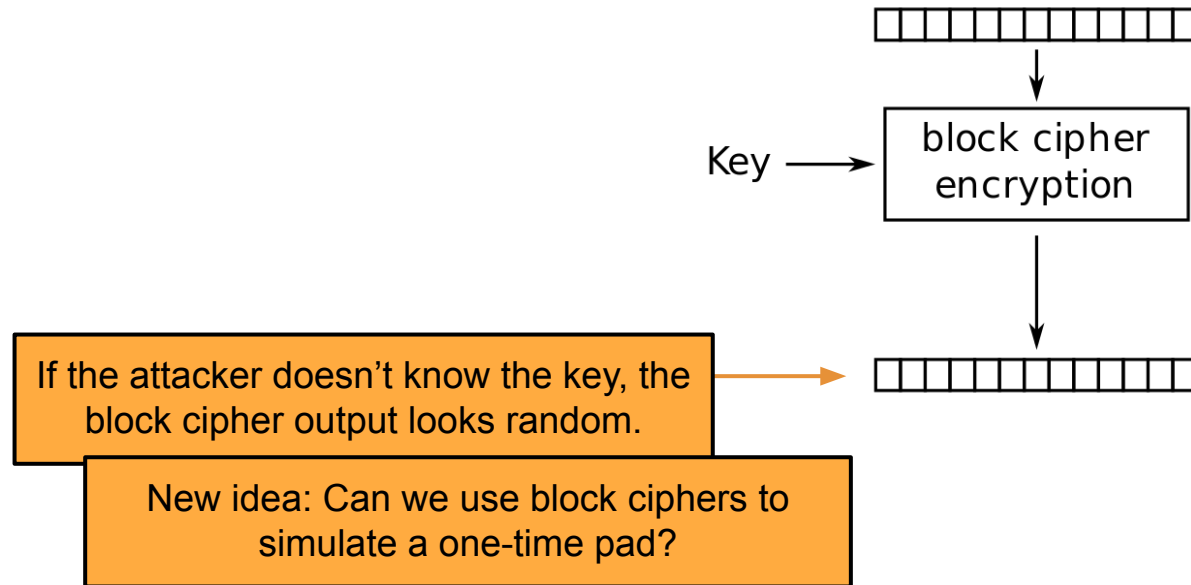
Computer Science 161

One-time pads are secure if we never reuse the key.




CTR Mode Scratchpad: Let's design it together

Computer Science 161



CTR Mode Scratchpad: Let's design it together

Computer Science 161






CTR Mode Scratchpad: Let's design it together

Computer Science 161

Idea: Use this random-looking output as a one-time pad!




Remember one-time pads: XOR the pad with plaintext to get ciphertext



CTR Mode Scratchpad: Let's design it together

Computer Science 161

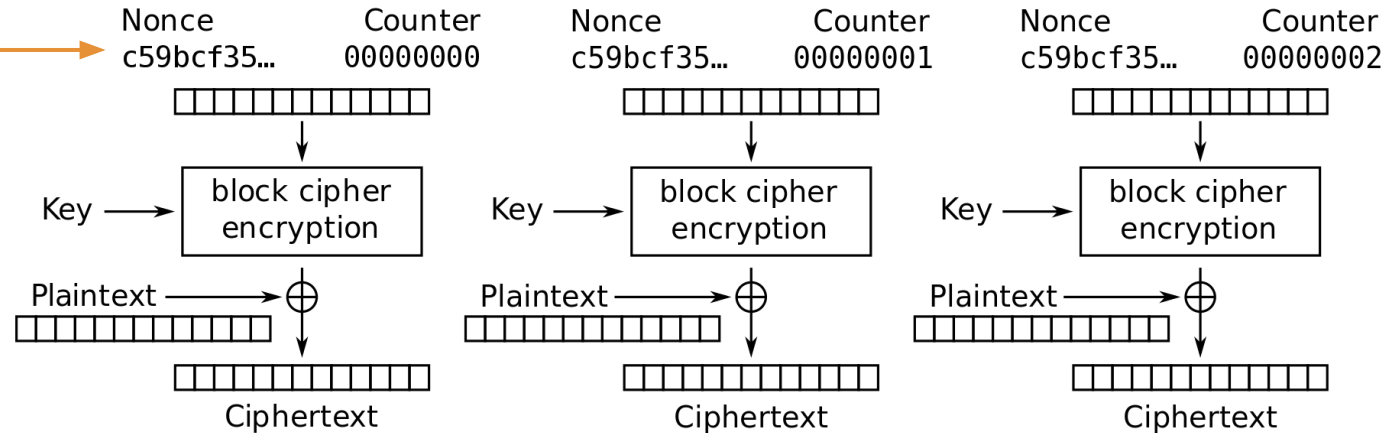
What do we use as input to the block cipher?



CTR Mode Scratchpad: Let's design it together

Computer Science 161

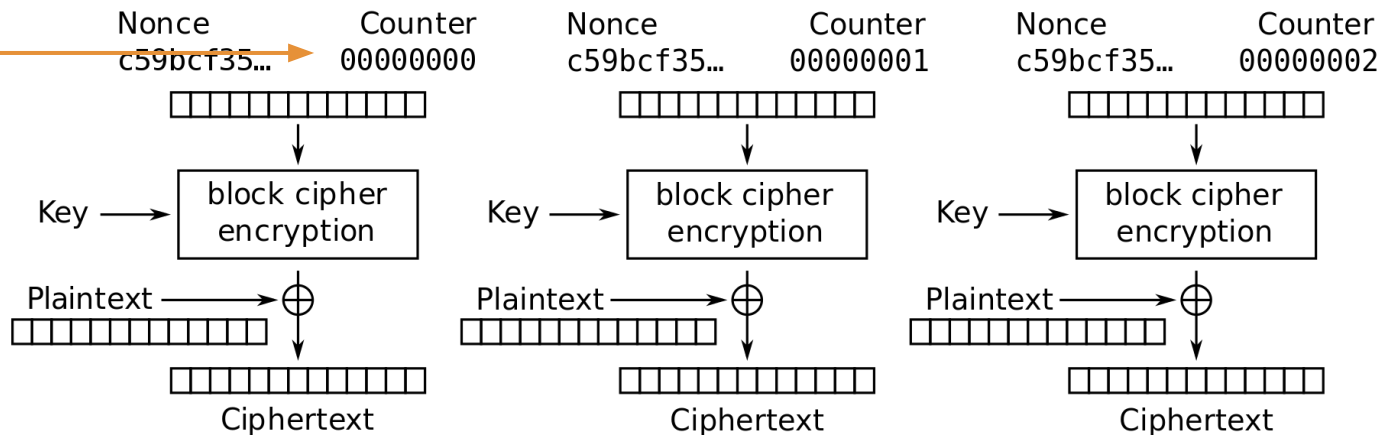
IND-CPA schemes need randomness, so let's put a random **nonce** here!



CTR Mode Scratchpad: Let's design it together

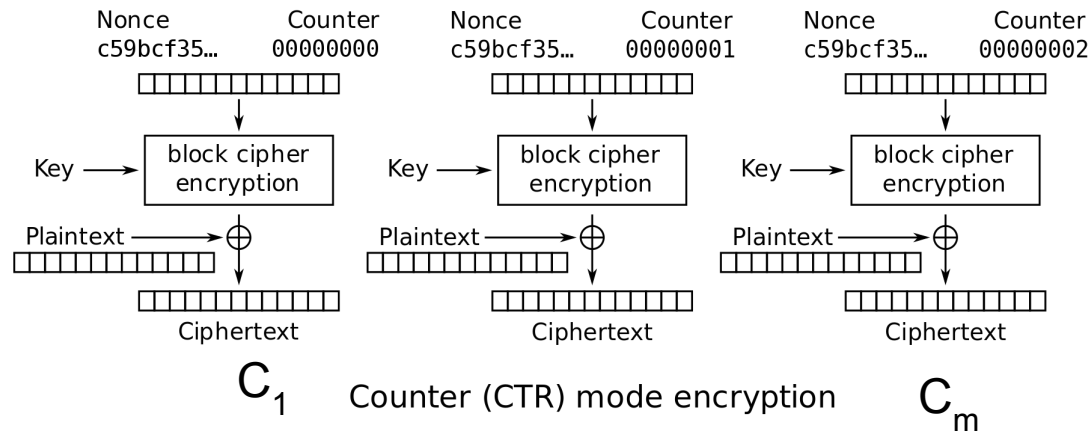
Computer Science 161

The **counter** increments per block to ensure each block cipher output is different.



CTR (Counter) Mode

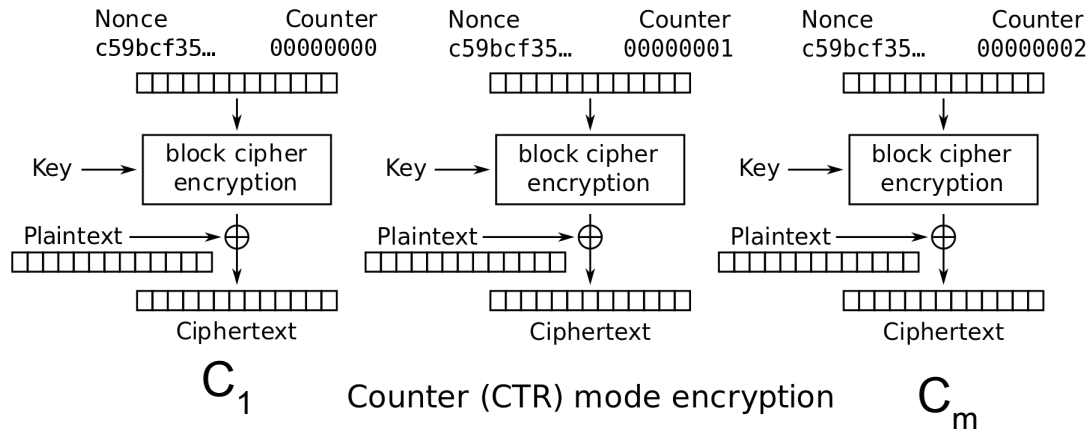
- Note: the random value is named the nonce here, but the idea is the same as the IV in CBC mode
- Overall ciphertext is (Nonce, C_1 , ..., C_m)



CTR Mode


- $\text{Enc}(K, M)$:
 - Split M in plaintext blocks $P_1 \dots P_m$ (each of block size n)
 - Choose random nonce
 - Compute and output $(\text{Nonce}, C_1, \dots, C_m)$

How do you decrypt?



CTR Mode: Decryption

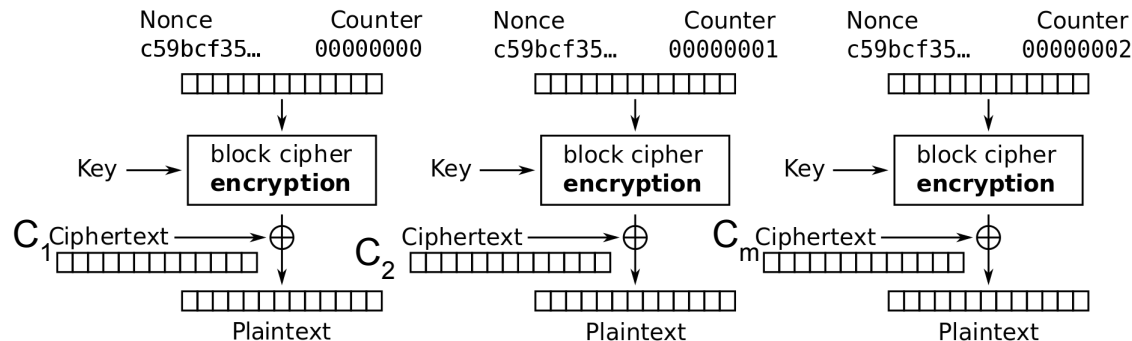
- Recall one-time pad: XOR with ciphertext to get plaintext
- Note: we are only using block cipher encryption, not decryption



Counter (CTR) mode decryption

CTR Mode: Decryption


- $\text{Dec}(K, C)$:
 - Parse C into (nonce, C_1, \dots, C_m)
 - Compute P_i by XORing C_i with output of E_k on nonce and counter
 - Concatenate resulting plaintexts and output $M = P_1 \dots P_m$




Counter (CTR) mode decryption

CTR Mode: Efficiency

- Can encryption be parallelized?
 - Yes
- Can decryption be parallelized?
 - Yes




Counter (CTR) mode encryption




Counter (CTR) mode decryption

CTR Mode: Padding

- Do we need to pad messages?
 - No! We can just cut off the parts of the XOR that are longer than the message.



Counter (CTR) mode encryption



Counter (CTR) mode decryption

CTR Mode: Security

- AES-CTR is IND-CPA secure. With what assumption?
- The nonce should never be reused (random generation helps here)
 - And in general less than $2^{n/2}$ blocks are encrypted
- What happens if you reuse the nonce?
- Equivalent to reusing a key in a one-time pad
 - Recall: Key reuse in a one-time pad is catastrophic: usually leaks enough information for an attacker to deduce the entire plaintext

CTR Mode: Penguin


Computer Science 161



Original image

CTR Mode: Penguin

Computer Science 161



Encrypted with CTR, with random nonces

IVs and Nonces

- **Initialization vector (IV):** A random, but public, one-use value to introduce randomness into the algorithm
 - For CTR mode, we say that you use a **nonce** (number used once), since the value has to be unique, not necessarily random.
 - In this class, we use IV and nonce interchangeably
- **Never reuse IVs**
 - In some algorithms, IV/nonce reuse leaks limited information (e.g. CBC)
 - In some algorithms, IV/nonce reuse leads to catastrophic failure (e.g. CTR)

IVs and Nonces

- Thinking about the consequences of IV/nonce reuse is hard
- What if the IV/nonce is not reused, but the attacker can predict future values?
 - Now you have to think about more attacks
 - We'll analyze this more in discussion: it really depends on the encryption function
- **Solution: Randomly generate a new IV/nonce for every encryption**
 - If the nonce is 128 bits or longer, the probability of generating the same IV/nonce twice is astronomically small (basically 0)
 - Now you don't have to think about IV/nonce reuse attacks!

The summer 2020 CS 61A exam mistake

Computer Science 161

- The TAs used a Python library for AES
 - A bad library for other reasons besides this example
- When they invoked CTR mode encryption, they didn't specify an IV
 - Assumption: the crypto library would add a random IV for them
 - Reality: the crypto library defaulted to IV = 0 every time
- The same IV was used to encrypt multiple exam questions
- All security was lost!
 - Any CS 161 student could have seen the exam beforehand
- **Takeaway:** Do not reuse IVs
- **Takeaway:** Real world cryptosystems are hard. Do *not* implement your own cryptosystems (without proper training beyond this class).

Comparing Modes of Operation

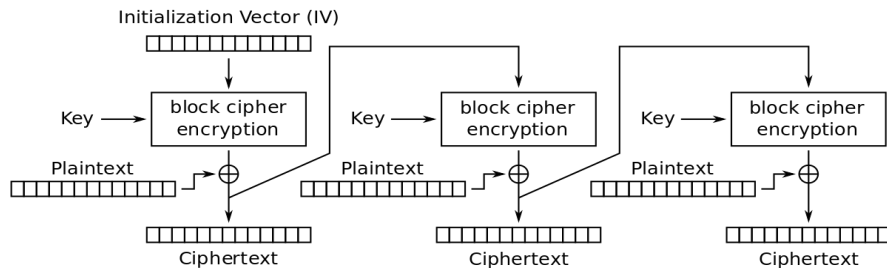
- If you need high performance, which mode is better?
 - CTR mode, because you can parallelize both encryption and decryption
- If you're paranoid about security, which mode is better?
 - CBC mode is better
- Theoretically, CBC and CTR mode are equally secure if used properly
 - However, if used improperly (IV/nonce reuse), CBC only leaks partial information, and CTR fails catastrophically
 - Consider human factors: Systems should be as secure as possible even when implemented *incorrectly*
 - IV failures on CTR mode have resulted in multiple real-world security incidents!

Other Modes of Operation

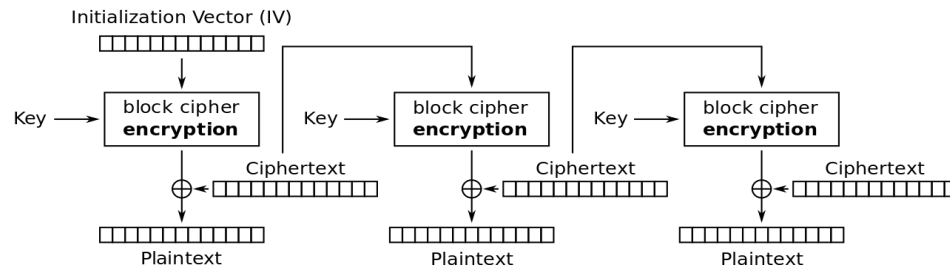
- Other modes exist besides CBC and CTR
- Trade-offs:
 - Do we need to pad messages?
 - How robust is the scheme if we use it incorrectly?
 - Can we parallelize encryption/decryption?

CFB Mode

- Also IND-CPA
- Try to analyze the trade-offs yourself:
 - Do we need to pad messages?
 - How robust is the scheme if we use it incorrectly?
 - Can we parallelize encryption/decryption?



Cipher Feedback (CFB) mode encryption




Cipher Feedback (CFB) mode decryption

CFB Mode

- Try to analyze the trade-offs yourself:
 - Do we need to pad messages?
 - No
 - How robust is the scheme if we use it incorrectly?
 - Similar effects as CBC mode, but a bit worse if you reuse the IV
 - Can we parallelize encryption/decryption?
 - Only decryption is parallelizable

Lack of Integrity and Authenticity

- Block ciphers are designed for *confidentiality* (IND-CPA)
- If an attacker tampers with the ciphertext, we are not guaranteed to detect it
- Remember Mallory: An *active* manipulator who wants to tamper with the message



Lack of Integrity and Authenticity

- Consider CTR mode
- What if Mallory tampers with the ciphertext using XOR?

	P	a	y		M	a	l		\$	1	0	0
M	0x50	0x61	0x79	0x20	0x4d	0x61	0x6c	0x20	0x24	0x31	0x30	0x30
\oplus												
$E_K(i)$	0x8a	0xe3	0x5e	0xcf	0x3b	0x40	0x46	0x57	0xb8	0x69	0xd2	0x96
$=$												
C	0xda	0x82	0x27	0xef	0x76	0x21	0x2a	0x77	0x9c	0x58	0xe2	0xa6

Lack of Integrity and Authenticity

- Suppose Mallory knows the message M
- How can Mallory change the M to say **Pay Mal \$900?**

M	P	a	y		M	a	l		\$	1	0	0
	0x50	0x61	0x79	0x20	0x4d	0x61	0x6c	0x20	0x24	0x31	0x30	0x30
\oplus												
$E_K(i)$	0x8a	0xe3	0x5e	0xcf	0x3b	0x40	0x46	0x57	0xb8	0x69	0xd2	0x96
$=$												
C	0xda	0x82	0x27	0xef	0x76	0x21	0x2a	0x77	0x9c	0x58	0xe2	0xa6

Lack of Integrity and Authenticity

$C_i = M_i \oplus \text{Pad}_i$	$0\mathbf{x58} = 0\mathbf{x31} \oplus \text{Pad}_i$	Definition of CTR
$\text{Pad}_i = M_i \oplus C_i$	$\text{Pad}_i = 0\mathbf{x58} \oplus 0\mathbf{x31}$ $= 0\mathbf{x69}$	Solve for the i th byte of the pad
$C'_i = M'_i \oplus \text{Pad}_i$	$C'_i = 0\mathbf{x39} \oplus 0\mathbf{x69}$ $= 0\mathbf{x50}$	Compute the changed i th byte

C	0xda	0x82	0x27	0xef	0x76	0x21	0x2a	0x77	0x9c	0x58	0xe2	0xa6
C'	0xda	0x82	0x27	0xef	0x76	0x21	0x2a	0x77	0x9c	0x50	0xe2	0xa6

Lack of Integrity and Authenticity

- What happens when we decrypt C' ?
 - The message looks like “Pay Mal \$900” now!
 - Note: Mallory didn’t have to know the key; no integrity or authenticity for CTR mode!

C'	0xda	0x82	0x27	0xef	0x76	0x21	0x2a	0x77	0x9c	0x50	0xe2	0xa6
------	------	------	------	------	------	------	------	------	------	------	------	------

\oplus

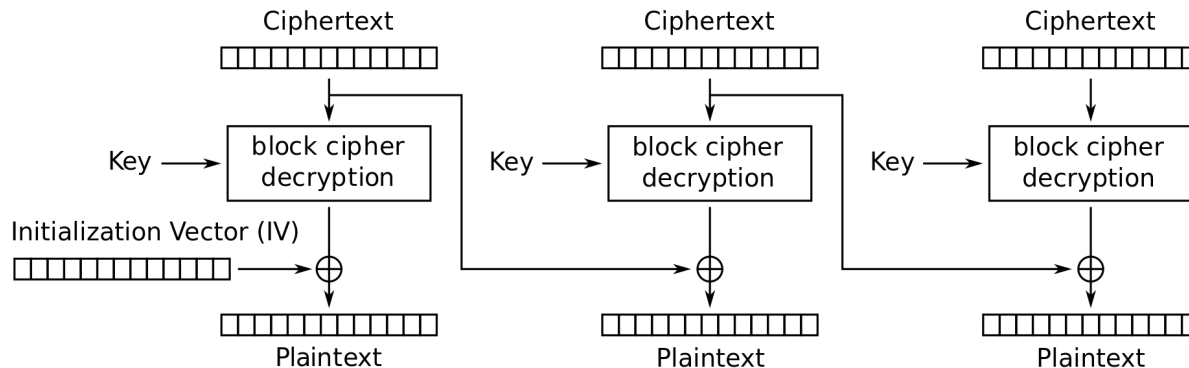
$E_K(i)$	0x8a	0xe3	0x5e	0xcf	0x3b	0x40	0x46	0x57	0xb8	0x69	0xd2	0x96
----------	------	------	------	------	------	------	------	------	------	------	------	------

=

P'	0x50	0x61	0x79	0x20	0x4d	0x61	0x6c	0x20	0x24	0x39	0x30	0x30
	P	a	y		M	a	l		\$	9	0	0

Lack of Integrity and Authenticity

- What about CBC?
 - Altering a bit of the ciphertext causes some blocks to become random gibberish
 - However, Bob doesn't know that Alice did not send random gibberish, so it still does *not* provide integrity or authenticity



Cipher Block Chaining (CBC) mode decryption

Today: Cryptography Hashes and MACs

Computer Science 161

- Hashing

- Definition
- Security: one-way, second preimage resistant, collision resistant
- Examples
- Length extension attacks
- Application: Lowest-hash scheme
- Do hashes provide integrity?

- MACs

- Definition
- Security: unforgeability
- Example: HMAC
- Do MACs provide integrity?

- Authenticated Encryption

- Definition
- Key Reuse
- MAC-then-Encrypt or Encrypt-then-MAC?
- AEAD Encryption Modes

Cryptographic Hashes



Textbook Chapter 7.1–7.3

Cryptography Roadmap

Computer Science 161

	Symmetric-key	Asymmetric-key
Confidentiality	<ul style="list-style-type: none">● One-time pads● Block ciphers with chaining modes (e.g. AES-CBC)	<ul style="list-style-type: none">● RSA encryption● ElGamal encryption
Integrity, Authentication	<ul style="list-style-type: none">● MACs (e.g. HMAC)	<ul style="list-style-type: none">● Digital signatures (e.g. RSA signatures)

- Hash functions
- Pseudorandom number generators
- Public key exchange (e.g. Diffie-Hellman)
- Key management (certificates)
- Password management

Cryptographic Hash Function: Definition

- Hash function: $H(M)$
 - Input: *Arbitrary* length message M
 - Output: *Fixed* length, n -bit hash
 - Sometimes written as $\{0, 1\}^* \rightarrow \{0, 1\}^n$
- Properties
 - **Correctness:** Deterministic
 - Hashing the same input always produces the same output
 - **Efficiency:** Efficient to compute
 - **Security:** One-way-ness (“preimage resistance”)
 - **Security:** Collision-resistance
 - **Security:** Random/unpredictability, no predictable patterns for how changing the input affects the output
 - Changing 1 bit in the input causes the output to be completely different
 - Also called “random oracle” assumption

Hash Function: Intuition

- A hash function provides a fixed-length “fingerprint” over a sequence of bits
- Example: Document comparison
 - If Alice and Bob both have a 1 GB document, they can both compute a hash over the document and (securely) communicate the hashes to each other
 - If the hashes are the same, the files must be the same, since they have the same “fingerprint”
 - If the hashes are different, the files must be different

Hash Function: One-way-ness or Preimage Resistance

- **Informal:** Given an output y , it is infeasible to find *any* input x such that $H(x) = y$
- **More formally:** For all polynomial time adversary,
Pr[x chosen randomly from plaintext space; $y = H(x)$:
Adv(y) outputs x' s.t. $H(x') = y$] is negligible
- Intuition: Here's an output. Can you find an input that hashes to this output?
 - Note: The adversary just needs to find *any* input, not necessarily the input that was actually used to generate the hash
- Example: Is $H(x) = 1$ one-way?
 - No, because given output 1, an attacker can return any number x
- Example: Is $H(\text{a cow}) = \text{a burger}$ one-way?
 - Most likely because you cannot come up with a cow that creates the exact burger

Hash Function: Collision Resistance

- **Collision:** Two different inputs with the same output
 - $x \neq x'$ and $H(x) = H(x')$
 - Can we design a hash function with no collisions?
 - No, because there are more inputs than outputs (pigeonhole principle)
 - However, we want to make finding collisions *infeasible* for an attacker
- **Collision resistance:** It is infeasible to (i.e. no polynomial time attacker can) find any pair of inputs $x' \neq x$ such that $H(x) = H(x')$
- Intuition: Can you find *any* two inputs that collide with the same hash output for *any* output?

Hash Function: Collision Resistance


Computer Science 161

- **Birthday attack:** Finding a collision on an n -bit output requires only $2^{n/2}$ tries on average
 - This is why a group of 23 people are >50% likely to have at least one birthday in common



Hash Function: Examples

- MD5
 - Output: 128 bits
 - Security: Completely broken
- SHA-1
 - Output: 160 bits
 - Security: Completely broken in 2017
 - Was known to be weak before 2017, but still used sometimes
- SHA-2
 - Output: 256, 384, or 512 bits (sometimes labeled SHA-256, SHA-384, SHA-512)
 - Not currently broken, but some variants are vulnerable to a length extension attack
 - Current standard
- SHA-3 (Keccak)
 - Output: 256, 384, or 512 bits
 - Current standard (not meant to replace SHA-2, just a different construction)



A GIF that displays its own MD5 hash

Length Extension Attacks

- **Length extension attack:** Given $H(x)$ and the length of x , but not x , an attacker can create $H(x || m)$ for any m of the attacker's choosing
 - Note: This doesn't violate any property of hash functions but is undesirable in some circumstances
- SHA-256 (256-bit version of SHA-2) is vulnerable
- SHA-3 is not vulnerable

Do hashes provide integrity?

- It depends on your threat model
- Scenario
 - Mozilla publishes a new version of Firefox on some download servers
 - Alice downloads the program binary
 - How can she be sure that nobody tampered with the program?
- Idea: use cryptographic hashes
 - Mozilla hashes the program binary and publishes the hash on its website
 - Alice hashes the binary she downloaded and checks that it matches the hash on the website
 - If Alice downloaded a malicious program, the hash would not match (tampering detected!)
 - An attacker can't create a malicious program with the same hash (collision resistance)
- Threat model: We assume the attacker cannot modify the hash on the website
 - We have integrity, as long as we can communicate the hash securely

Do hashes provide integrity?

- It depends on your threat model
- Scenario
 - Alice and Bob want to communicate over an insecure channel
 - Mallory might tamper with messages
- Idea: Use cryptographic hashes
 - Alice sends her message with a cryptographic hash over the channel
 - Bob receives the message and computes a hash on the message
 - Bob checks that the hash he computed matches the hash sent by Alice
- Threat model: Mallory can modify the message *and the hash*
 - No integrity!

Do hashes provide integrity?

- It depends on your threat model
- If the attacker can modify the hash, hashes don't provide integrity
- Main issue: Hashes are *unkeyed* functions
 - There is no secret key being used as input, so any attacker can compute a hash on any value
- Next: Use hashes to design schemes that provide integrity

Message Authentication Codes (MACs)



Textbook Chapter 8.1–8.3 & 8.5–8.6

Cryptography Roadmap

Computer Science 161

	Symmetric-key	Asymmetric-key
Confidentiality	<ul style="list-style-type: none">● One-time pads● Block ciphers with chaining modes (e.g. AES-CBC)	<ul style="list-style-type: none">● RSA encryption● ElGamal encryption
Integrity, Authentication	<ul style="list-style-type: none">● MACs (e.g. HMAC)	<ul style="list-style-type: none">● Digital signatures (e.g. RSA signatures)

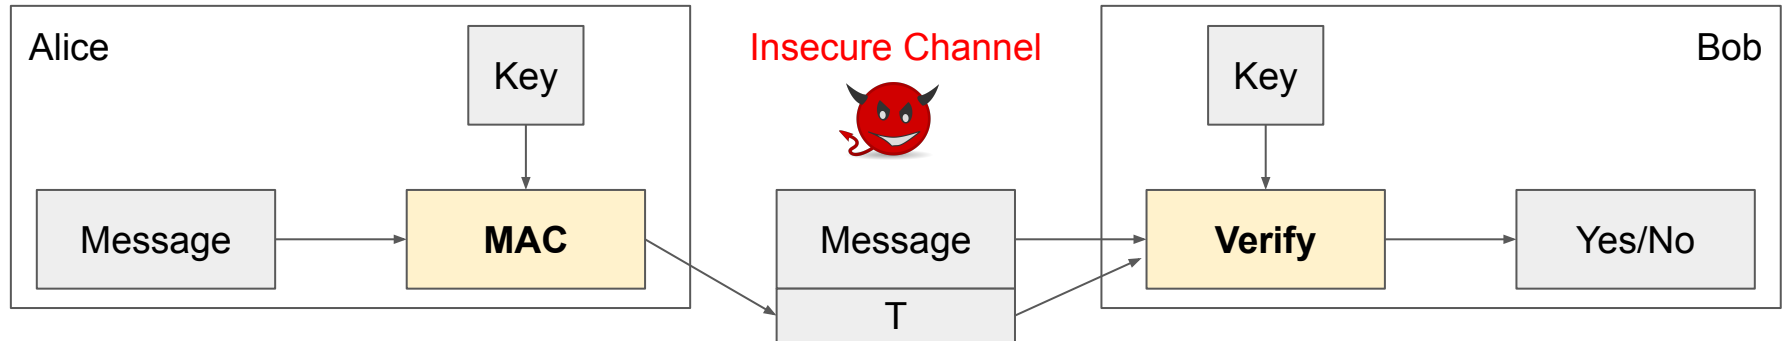
- Hash functions
- Pseudorandom number generators
- Public key exchange (e.g. Diffie-Hellman)
- Key management (certificates)
- Password management

How to Provide Integrity

- Reminder: We're still in the symmetric-key setting
 - Assume that Alice and Bob share a secret key, and attackers don't know the key
- We want to attach some piece of information to *prove* that someone with the key sent this message
 - This piece of information can only be generated by someone with the key

MACs: Usage

- Alice wants to send M to Bob, but doesn't want Mallory to tamper with it
- Alice sends M and $T = \text{MAC}(K, M)$ to Bob
- Bob recomputes $\text{MAC}(K, M)$ and checks that it matches T
- If the MACs match, Bob is confident the message has not been tampered with (integrity)



MACs: Definition

- Two parts:
 - $\text{KeyGen}() \rightarrow K$: Generate a key K
 - $\text{MAC}(K, M) \rightarrow T$: Generate a tag T for the message M using key K
 - Inputs: A secret key and an arbitrary-length message
 - Output: A fixed-length **tag** on the message
- Properties
 - **Correctness**: Determinism
 - Note: Some more complicated MAC schemes have an additional $\text{Verify}(K, M, T)$ function that don't require determinism, but this is out of scope
 - **Efficiency**: Computing a MAC should be efficient
 - **Security**: EU-CPA (existentially unforgeable under chosen plaintext attack)

Defining Integrity: EU-CPA

- A secure MAC is **existentially unforgeable**: without the key, an attacker cannot create a valid tag on a message
 - Mallory cannot generate $\text{MAC}(K, M')$ without K
 - Mallory cannot find any $M' \neq M$ such that $\text{MAC}(K, M') = \text{MAC}(K, M)$
- Formally defined by a security game: existential unforgeability under chosen-plaintext attack, or EU-CPA
- MACs should be unforgeable under chosen plaintext attack
 - Intuition: Like IND-CPA, but for integrity and authenticity
 - Even if Mallory can trick Alice into creating MACs for messages that Mallory chooses, Mallory cannot create a valid MAC on a message that she hasn't seen before

Defining Integrity: EU-CPA

1. Mallory may send messages to Alice and receive their tags
2. Eventually, Mallory creates a message-tag pair (M', T')
 - M' cannot be a message that Mallory requested earlier
 - If T' is a valid tag for M' , then Mallory wins. Otherwise, she loses.
3. A scheme is EU-CPA secure if for *all* polynomial time adversaries, the probability of winning is 0 or negligible

