

MACs, PRNGs and Diffie-Hellman Key Exchange

CS 161 Spring 2023 - Lecture 9

Announcements

- **HW 3 is released and due this Friday, February 16th at 11:59 PT**
- **The midterm is on Thursday, February 29th from 7:00-9:00 PM PT.**
 - If you would like to request an alternate exam time or remote exam, or have DSP accommodations or any special requests, please **fill out the [Exam Logistics Form](#) by Monday, February 19, 11:59 PM PT.**
 - We'll announce more information regarding review sessions, past exams, etc next week.

Last Time: Hashes

Computer Science 161

- Map arbitrary-length input to fixed-length output
- Output is deterministic and unpredictable
- Security properties
 - One way: Given an output y , it is infeasible to find any input x such that $H(x) = y$.
 - Collision resistant: It is infeasible to find another any pair of inputs $x' \neq x$ such that $H(x) = H(x')$.
- Some hashes are vulnerable to length extension attacks
- Hashes don't provide integrity (unless you can publish the hash securely)

Message Authentication Codes (MACs)



Textbook Chapter 8.1–8.3 & 8.5–8.6

Cryptography Roadmap

Computer Science 161

| | Symmetric-key | Asymmetric-key |
|---------------------------|--|---|
| Confidentiality | <ul style="list-style-type: none">● One-time pads● Block ciphers with chaining modes (e.g. AES-CBC) | <ul style="list-style-type: none">● RSA encryption● ElGamal encryption |
| Integrity, Authentication | <ul style="list-style-type: none">● MACs (e.g. HMAC) | <ul style="list-style-type: none">● Digital signatures (e.g. RSA signatures) |

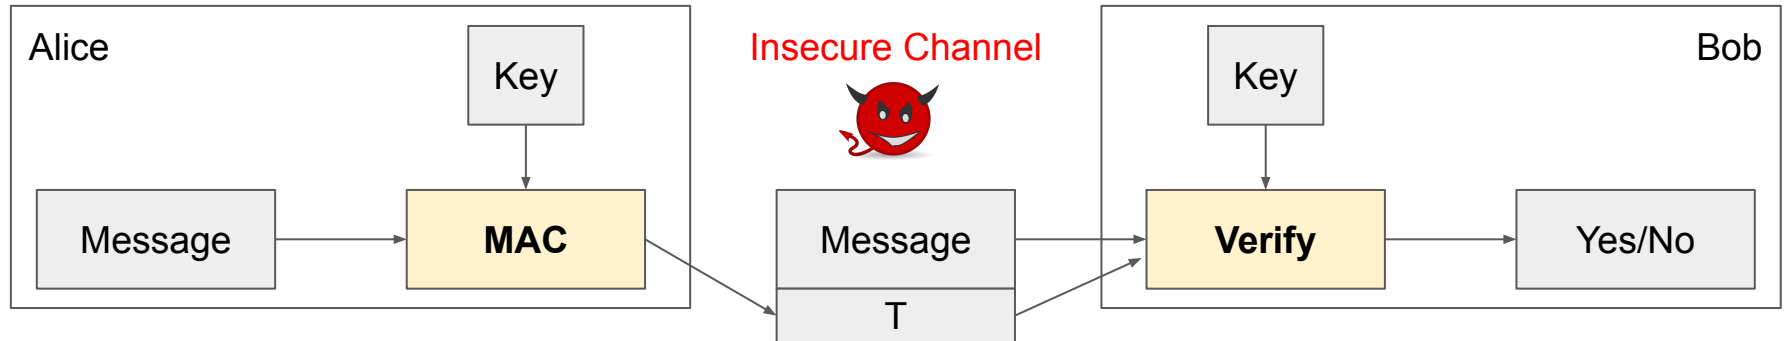
- Hash functions
- Pseudorandom number generators
- Public key exchange (e.g. Diffie-Hellman)
- Key management (certificates)
- Password management

How to Provide Integrity

- Reminder: We're still in the symmetric-key setting
 - Assume that Alice and Bob share a secret key, and attackers don't know the key
- We want to attach some piece of information to *convince* Bob that Alice sent this message even if Mallory is intercepting the message on the network, or to *detect* if Mallory tampered with the message
 - This piece of information can only be generated by someone with the key

MACs: Usage

- Alice wants to send M to Bob, but doesn't want **Mallory** to tamper with it
- Alice sends M and $T = \text{MAC}(K, M)$ to Bob
- Bob recomputes $\text{MAC}(K, M)$ and checks that it matches T
- If the MACs match, Bob is confident the message has not been tampered with (integrity)



MACs: Definition

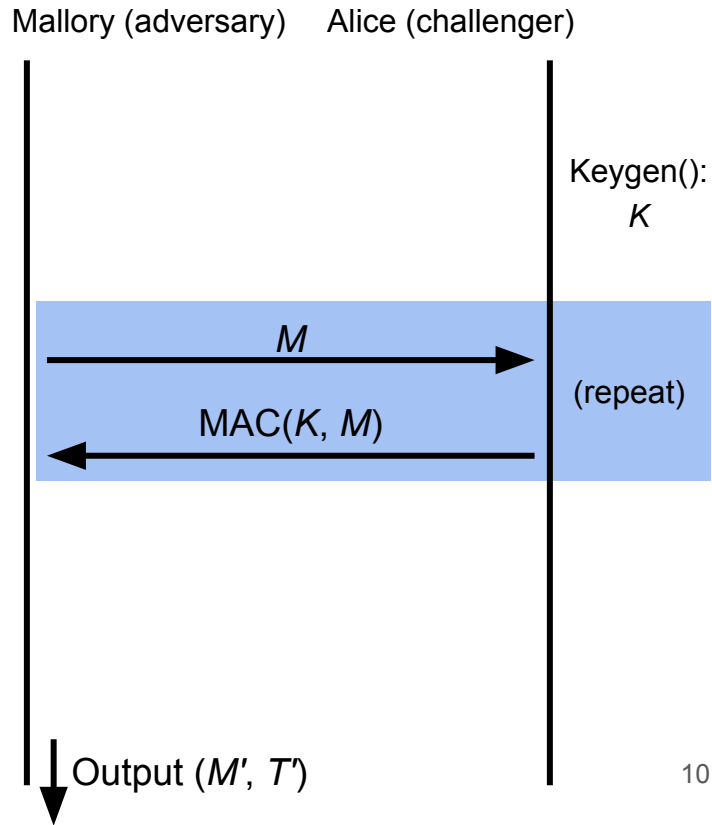
- Two parts:
 - $\text{KeyGen}() \rightarrow K$: Generate a key K
 - $\text{MAC}(K, M) \rightarrow T$: Generate a tag T for the message M using key K
 - Inputs: A secret key and an arbitrary-length message
 - Output: A fixed-length **tag** on the message
- Properties
 - **Correctness**: Determinism
 - Note: Some more complicated MAC schemes have an additional $\text{Verify}(K, M, T)$ function that don't require determinism, but this is out of scope
 - **Efficiency**: Computing a MAC should be efficient
 - **Security**: EU-CPA (existentially unforgeable under chosen plaintext attack)

Defining Integrity: EU-CPA

- A secure MAC is **existentially unforgeable**: without the key, an attacker cannot create a valid tag on a message
 - Mallory cannot generate $\text{MAC}(K, M')$ without K
 - Mallory cannot find any $M' \neq M$ such that $\text{MAC}(K, M') = \text{MAC}(K, M)$
- Formally defined by a security game: existential unforgeability under chosen-plaintext attack, or EU-CPA
- MACs should be unforgeable under chosen plaintext attack
 - Intuition: Like IND-CPA, but for integrity and authenticity
 - Even if Mallory can trick Alice into creating MACs for messages that Mallory chooses, Mallory cannot create a valid MAC on a message that she hasn't seen before

Defining Integrity: EU-CPA

1. Mallory may send messages to Alice and receive their tags
2. Eventually, Mallory creates a message-tag pair (M', T')
 - M' cannot be a message that Mallory requested earlier
 - If T' is a valid tag for M' , then Mallory wins. Otherwise, she loses.
3. A scheme is EU-CPA secure if for *all* polynomial time adversaries, the probability of winning is 0 or negligible



Example: NMAC

- Can we use secure cryptographic hashes to build a secure MAC?
 - Intuition: Hash output is unpredictable and looks random, so let's hash the key and the message together
- KeyGen():
 - Output two random, n -bit keys K_1 and K_2 , where n is the length of the hash output
- NMAC(K_1, K_2, M):
 - Output $H(K_1 \parallel H(K_2 \parallel M))$
- NMAC is EU-CPA secure if the two keys are different
 - Provably secure if the underlying hash function is secure
- Intuition: Using two hashes prevents a length extension attack
 - Otherwise, an attacker who sees a tag for M could generate a tag for $M \parallel M'$

Example: HMAC

- Issues with NMAC:
 - Recall: $\text{NMAC}(K_1, K_2, M) = H(K_1 \parallel H(K_2 \parallel M))$
 - We need two different keys
 - NMAC requires the keys to be the same length as the hash output (n bits)
- $\text{HMAC}(K, M)$:
 - Compute K' as a version of K that is the length of the hash output
 - If K is too short, pad K with 0's to make it n bits (be careful with keys that are too short and lack randomness)
 - If K is too long, hash it so it's n bits
 - Output $H(K' \oplus \text{opad} \parallel H(K' \oplus \text{ipad} \parallel M))$

Example: HMAC

- HMAC(K , M):
 - Compute K' as a version of K that is the length of the hash output
 - If K is too short, pad K with 0's to make it n bits (be careful with keys that are too short and lack randomness)
 - If K is too long, hash it so it's n bits
 - Output $H((K' \oplus \text{opad}) \parallel H((K' \oplus \text{ipad}) \parallel M))$
- Use K' to derive two different keys
 - *opad* (outer pad) is the hard-coded byte `0x5c` repeated until it's the same length as K'
 - *ipad* (inner pad) is the hard-coded byte `0x36` repeated until it's the same length as K'
 - As long as *opad* and *ipad* are different, you'll get two different keys
 - For paranoia, the designers chose two very different bit patterns, even though they theoretically need only differ in one bit

HMAC Properties

Computer Science 161

- $\text{HMAC}(K, M) = H((K \oplus \text{opad}) \parallel H((K \oplus \text{ipad}) \parallel M))$
- HMAC is a hash function, so it has the properties of the underlying hash too
 - It is collision resistant
 - Given $\text{HMAC}(K, M)$ and K , an attacker can't learn M
 - If the underlying hash is secure, HMAC doesn't reveal M , but it is still deterministic
- You can't verify a tag T if you don't have K
 - The attacker can't brute-force the message M without knowing K

Do MACs provide integrity?

- Do MACs provide integrity?
 - Yes. An attacker cannot tamper with the message without being detected
- Do MACs provide authenticity?
 - It depends on your threat model
 - If a message has a valid MAC, you can be sure it came from *someone with the secret key*, but you can't narrow it down to one person
 - If only two people have the secret key, MACs provide authenticity: it has a valid MAC, and it's not from me, so it must be from the other person
- Do MACs provide confidentiality?
 - MACs are deterministic \Rightarrow No IND-CPA security
 - MACs in general have no confidentiality guarantees; they can leak information about the message

MACs: Summary

- Inputs: a secret key and a message
- Output: a tag on the message
- A secure MAC is unforgeable: Even if Mallory can trick Alice into creating MACs for messages that Mallory chooses, Mallory cannot create a valid MAC on a message that she hasn't seen before
 - Example: $\text{HMAC}(K, M) = H((K' \oplus \text{opad}) \parallel H((K' \oplus \text{ipad}) \parallel M))$
- MACs do not provide confidentiality

Authenticated Encryption



Textbook Chapter 8.7 & 8.8

Cryptography Roadmap

Computer Science 161

| | Symmetric-key | Asymmetric-key |
|---------------------------|--|---|
| Confidentiality | <ul style="list-style-type: none">● One-time pads● Block ciphers with chaining modes (e.g. AES-CBC) | <ul style="list-style-type: none">● RSA encryption● ElGamal encryption |
| Integrity, Authentication | <ul style="list-style-type: none">● MACs (e.g. HMAC) | <ul style="list-style-type: none">● Digital signatures (e.g. RSA signatures) |

- Hash functions
- Pseudorandom number generators
- Public key exchange (e.g. Diffie-Hellman)
- Key management (certificates)
- Password management

Authenticated Encryption: Definition

- **Authenticated encryption (AE):** A scheme that simultaneously guarantees confidentiality and integrity (and authenticity, depending on your threat model) on a message
- Two ways of achieving authenticated encryption:
 - Combine schemes that provide confidentiality with schemes that provide integrity
 - Use a scheme that is designed to provide confidentiality and integrity

Combining Schemes: Let's design it together

- You can use:
 - An IND-CPA encryption scheme (e.g. AES-CBC): $\text{Enc}(K, M)$ and $\text{Dec}(K, M)$
 - An unforgeable MAC scheme (e.g. HMAC): $\text{MAC}(K, M)$
- First attempt: Alice sends $\text{Enc}(K_1, M)$ and $\text{MAC}(K_2, M)$
 - Integrity? Yes, attacker can't tamper with the MAC
 - Confidentiality? No, the MAC is not IND-CPA secure
- Idea: Let's compute the MAC on the *ciphertext* instead of the plaintext:
 $\text{Enc}(K_1, M)$ and $\text{MAC}(K_2, \text{Enc}(K_1, M))$
 - Integrity? Yes, attacker can't tamper with the MAC
 - Confidentiality? Yes, the MAC might leak info about the ciphertext, but that's okay
- Idea: Let's encrypt the MAC too: $\text{Enc}(K_1, M \parallel \text{MAC}(K_2, M))$
 - Integrity? Yes, attacker can't tamper with the MAC
 - Confidentiality? Yes, everything is encrypted

MAC-then-Encrypt or Encrypt-then-MAC?

Computer Science 161

- MAC-then-encrypt
 - First compute $\text{MAC}(K_2, M)$
 - Then encrypt the message and the MAC together: $\text{Enc}(K_1, M \parallel \text{MAC}(K_2, M))$
- Encrypt-then-MAC
 - First compute $\text{Enc}(K_1, M)$
 - Then MAC the ciphertext: $\text{MAC}(K_2, \text{Enc}(K_1, M))$
- Which is better?
 - In theory, both are IND-CPA and EU-CPA secure if applied properly
 - MAC-then-encrypt has a downside: You don't know if tampering has occurred until after decrypting
 - Attacker can supply arbitrary tampered input, and you always have to decrypt it
 - Passing attacker-chosen input through the decryption function can cause side-channel leaks
- **Always use encrypt-then-MAC** because it's more robust to mistakes

Key Reuse

- **Key reuse problem:** Using the same key in two different use cases
 - Note: Using the same key multiple times for the same use (e.g. computing HMACs on different messages in the same context with the same key) is not key reuse problem
- Reusing keys can cause the underlying algorithms to interfere with each other and affect security guarantees
 - Example: If you use a block-cipher-based MAC algorithm and a block cipher chaining mode, the underlying block ciphers may no longer be secure
 - Thinking about these attacks is hard

Key Reuse

- Simplest solution: Do not reuse keys across schemes! One key per *scheme instance*.
 - Encrypt a piece of data, and MAC a piece of data?
 - Different use; different key
 - MAC one of Alice's messages to Bob and MAC one of Bob's messages to Alice?
 - Different use; different key
 - Encrypt one of Alice's files and encrypt another one of Alice's files?
 - It's *probably* fine to use the same key, but cryptographic design is tricky to get right!
 - Encrypt user metadata, encrypt file metadata, and encrypt file data?
 - You'll have to think about this in Project 2!

TLS 1.0 “Lucky 13” Attack

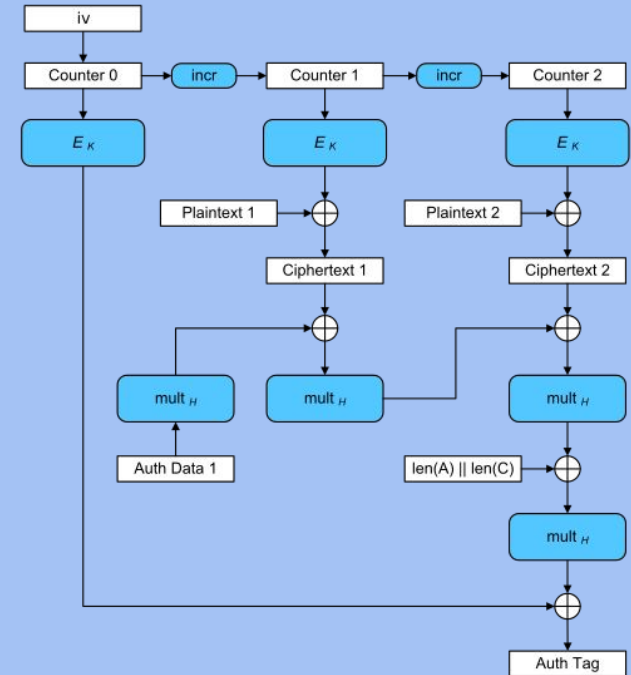
- TLS: A protocol for sending encrypted and authenticated messages over the Internet (we’ll study it more in the networking unit)
- TLS 1.0 uses MAC-then-encrypt: $\text{Enc}(K_1, M \parallel \text{MAC}(K_2, M))$
 - The encryption algorithm is AES-CBC
- The Lucky 13 attack abuses MAC-then-encrypt to read encrypted messages
 - Guess a byte of plaintext and change the ciphertext accordingly
 - The MAC will error, but the time it takes to error is different depending on if the guess is correct
 - Attacker measures how long it takes to error in order to learn information about plaintext
 - TLS will send the message again if the MAC errors, so the attacker can guess repeatedly
- Takeaways
 - Side channel attack: The algorithm is proved secure, but poor implementation made it vulnerable
 - Always encrypt-then-MAC

AEAD Encryption

- Second method for authenticated encryption: Use a scheme that is designed to provide confidentiality, integrity, and authenticity
- **Authenticated encryption with additional data (AEAD)**: An algorithm that provides both confidentiality and integrity over the plaintext and integrity over *additional data*
 - Additional data is usually context (e.g. memory address), so you can't change the context without breaking the MAC
- Great if used correctly: No more worrying about MAC-then-encrypt
 - If you use AEAD incorrectly, you lose *both* confidentiality and integrity/authentication
 - Example of correct usage: Using a crypto library with AEAD

AEAD Example: Galois Counter Mode (GCM)


- **Galois Counter Mode (GCM):** An AEAD block cipher mode of operation
- E_K is standard block cipher encryption
- mult_H is 128-bit multiplication over a special field (Galois multiplication)
 - Don't worry about the math



AEAD Example: Galois Counter Mode (GCM)

Computer Science 161


- Very fast mode of operation
 - Fully parallel encryption
 - Galois multiplication isn't parallelizable, but it's very fast
- Drawbacks
 - IV reuse leads to loss of confidentiality, integrity, and authentication
 - This wouldn't happen if you used AES-CTR and HMAC-SHA256
 - Implementing Galois is difficult and easy to screw up
- **Takeaway:** GCM provides integrity and confidentiality, but if you misuse it, it's even worse than CTR mode



Authenticated Encryption: Summary

- Authenticated encryption: A scheme that simultaneously guarantees confidentiality and integrity (and authenticity) on a message
- First approach: Combine schemes that provide confidentiality with schemes that provide integrity and authenticity
 - MAC-then-encrypt: $\text{Enc}(K_1, M \parallel \text{MAC}(K_2, M))$
 - Encrypt-then-MAC: $\text{Enc}(K_1, M) \parallel \text{MAC}(K_2, \text{Enc}(K_1, M))$
 - Always use Encrypt-then-MAC because it's more robust to mistakes
- Second approach: Use AEAD encryption modes designed to provide confidentiality, integrity, and authenticity
 - Drawback: Incorrectly using AEAD modes leads to losing *both* confidentiality and integrity/authentication

Pseudorandom Number Generators (PRNGs)



Textbook Chapter 9

Cryptography Roadmap

Computer Science 161

| | Symmetric-key | Asymmetric-key |
|---------------------------|--|---|
| Confidentiality | <ul style="list-style-type: none">● One-time pads● Block ciphers with chaining modes (e.g. AES-CBC) | <ul style="list-style-type: none">● RSA encryption● ElGamal encryption |
| Integrity, Authentication | <ul style="list-style-type: none">● MACs (e.g. HMAC) | <ul style="list-style-type: none">● Digital signatures (e.g. RSA signatures) |

- Hash functions
- Pseudorandom number generators
- Public key exchange (e.g. Diffie-Hellman)
- Key management (certificates)
- Password management

Randomness

- Randomness is essential for symmetric-key encryption
 - A random key
 - A random IV/nonce
 - Universally unique identifiers (we'll see this shortly)
 - We'll see more applications later
- If an attacker can predict a random number, things can catastrophically fail
- How do we securely generate random numbers?

Entropy

- In cryptography, “random” usually means “random and unpredictable”
- Scenario
 - You want to generate a secret bitstring that the attacker can't guess
 - You generate random bits by tossing a fair (50-50) coin
 - The outcomes of the fair coin are harder for the attacker to guess
- **Entropy: A measure of uncertainty**
 - In other words, a measure of how unpredictable the outcomes are
 - High entropy = unpredictable outcomes = desirable in cryptography
 - The uniform distribution has the highest entropy (every outcome equally likely, e.g. fair coin toss)
 - Usually measured in bits (so 3 bits of entropy = uniform, random distribution over 8 values)

Breaking Bitcoin Wallets

Computer Science 161

- What happens if we use a poor source of entropy?
- Bitcoin users use a randomly-generated private key to access their account (and money)
 - An attacker who learns the key can access the money
 - We'll learn more about Bitcoin later
- An “improvement” [sic] to the algorithm reduced the entropy used to generate the private keys
 - Any private key created with this “improvement” could be brute-forced

Improvements to RNG

committed on Dec 7, 2014 1 parent b0d5639

Showing 1 changed file with 26 additions and 28 deletions.

54 bitcoinjs-lib/src/jsbn/rng.js

```
@@ -8,15 +8,16 @@ var rng_state;
8      var rng_pool;
9      var rng_pptr;
10
11 - // Mix in a 32-bit integer into the pool
12 - function rng_seed_int(x) {
13 -     rng_pool[rng_pptr++] ^= x & 255;
14 -     rng_pool[rng_pptr++] ^= (x >> 8) & 255;
15 -     rng_pool[rng_pptr++] ^= (x >> 16) & 255;
16 -     rng_pool[rng_pptr++] ^= (x >> 24) & 255;
```

True Randomness

Computer Science 161

- To generate truly random numbers, we need a physical source of entropy
 - An unpredictable circuit on a CPU
 - Human activity measured at very fine time scales (e.g. the microsecond you pressed a key)
- Unbiased entropy usually requires combining multiple entropy sources
 - Goal: Total number of bits of entropy is the sum of all the input numbers of bits of entropy
 - Many poor sources + 1 good source = good entropy
- Issues with true randomness
 - It's expensive and slow to generate
 - Physical entropy sources are often biased



Exotic entropy source: Cloudflare has a wall of lava lamps that are recorded by an HD video camera that views the lamps through a rotating prism

Pseudorandom Number Generators (PRNGs)

- True randomness is expensive and biased
- **Pseudorandom number generator (PRNGs)**: An algorithm that uses a little bit of true randomness to generate a lot of random-looking output
 - Also called **deterministic random bit generators (DRBGs)**
- Usage
 - Generate some expensive true randomness (e.g. noisy circuit on your CPU)
 - Use the true randomness as input to the PRNG
 - Generate random-looking numbers quickly and cheaply with the PRNG
- PRNGs are deterministic: Output is generated according to a set algorithm
 - However, for an attacker who can't see the internal state, the output is *computationally indistinguishable* from true randomness

Stream Ciphers

Textbook Chapter 9.5

PRNG: Definition

- A PRNG has two functions:
 - PRNG.Seed(randomness): Initializes the internal state using the entropy
 - Input: Some truly random bits
 - PRNG.Generate(m): Generate m pseudorandom bits
 - Input: A number m
 - Output: m pseudorandom bits
 - Updates the internal state as needed

Properties

- **Correctness:** Deterministic
- **Efficiency:** Efficient to generate pseudorandom bits
- **Security:** Indistinguishability from random

PRNG: Security

- Can we design a PRNG that is truly random?
- A PRNG cannot be truly random
 - The output is deterministic given the initial seed
 - If the initial seed is s bits long, there are only 2^s possible output sequences
- A secure PRNG is computationally indistinguishable from random to an attacker
 - Game: Present an attacker with a truly random sequence and a sequence outputted from a secure PRNG
 - An attacker should not be able to determine which is which with probability $> \frac{1}{2} + \text{negl}$
- Equivalence: An attacker cannot predict future output of the PRNG

Insecure PRNGs: Breaking Slot Machines

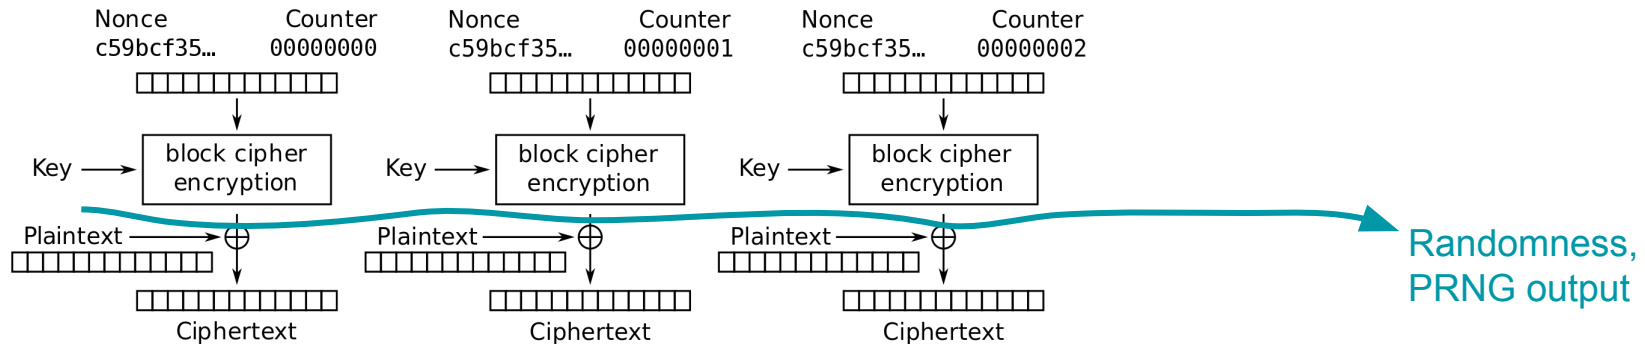
- What happens if PRNGs are used improperly?
- A casino in St. Louis experienced unusual bad “luck”
 - Suspicious players would hover over the lever and then spin at a specific time to win
- Vulnerability: Slot machines used predictable PRNGs
 - The PRNG output was based on the current time
- Strategy:
 - Use a smartphone to alert you to when to pull the lever for the best chance of winning
- Las Vegas was not affected by the vulnerability
 - Nevada slot machines must follow evaluation standards designed to address this sort of issue

Insecure PRNGs: OpenSSL PRNG bug

- What happens if we don't use enough entropy?
- Debian OpenSSL CVE-2008-0166
 - Debian: A Linux distribution
 - OpenSSL: A cryptographic library
 - In “cleaning up” OpenSSL (Debian “bug” #363516), the author “fixed” how OpenSSL seeds random numbers
 - The existing code caused Purify and Valgrind to complain about reading uninitialized memory
 - The cleanup caused the PRNG to only be seeded with the process ID
 - There are only 2^{15} (32,768) possible process IDs, so the PRNG only has 15 bits of entropy
- Easy to deduce private keys generated with the PRNG
 - Set the PRNG to every possible starting state and generate a few private/public key pairs
 - See if the matching public key is anywhere on the Internet

Example construction of PRNG

- Using block cipher in CTR mode:
- If you want m random bits, and a block cipher with E_k has n bits, apply the block cipher m/n times and concatenate the result:
- $\text{PRNG.Seed}(K \mid \text{IV});$
- $\text{Generate}(m) = E_k(\text{IV} \mid 1) \mid E_k(\text{IV} \mid 2) \mid E_k(\text{IV} \mid 3) \dots E_k(\text{IV} \mid \text{ceil}(m/n)),$
 - \mid is concatenation



Insecure PRNGs: Rust Rand_Core

- A Rust library has an interface for “secure” random number generators... but it isn’t actually secure!
- Example: ChaCha8Rng
 - A stream cipher PRNG
 - No reseed function: no way of adding extra entropy after the initial seed
 - Seed only takes 32 bits: no way to combine entropy
 - No rollback resistance
- None of the “secure” RNGs are cryptographically secure
 - None have a reseed function to add extra entropy
 - None take arbitrarily long seeds
- **Takeaway:** Always make sure you use a secure PRNG
 - Consider human factors? Use fail-safe defaults?

Insecure PRNGs: CVE-2019-16303

- Relevant if you wrote an app in JHipster before 2019
- Password reset functions
 - When you forget your password, receive an email with a special link to reset your password
 - The special link should contain a randomly-generated code (so attackers can't make their own link)
- Vulnerability: Bad PRNG
 - You can figure out the PRNG's internal state from the reset link
 - Request password reset links for other people's accounts
 - Predict the “random” reset link and take over any account you want!

Application: Universally Unique Identifiers (UUIDs)

- Scenario
 - You have a set of objects (e.g. files)
 - You need to assign a unique name to every object
 - Every name must be unique and unpredictable
- Solution: choose a random value
 - If you use enough randomness, the probability of generating the same random value twice are astronomically small (basically 0)
- Universally Unique Identifiers (UUIDs)
 - 128-bit unique values
 - To generate a new UUID, seed a secure PRNG properly, and generate a random value
 - Often written in hexadecimal: **00112233-4455-6677-8899-aabbccddeeff**
 - You'll work with UUIDs in Project 2

PRNGs: Summary

- True randomness requires sampling a physical process
 - Slow, expensive, and biased (low entropy)
- PRNG: An algorithm that uses a little bit of true randomness to generate a lot of random-looking output
 - Seed(entropy): Initialize internal state
 - Generate(n): Generate n bits of pseudorandom output
- Security: computationally indistinguishable from truly random bits
- Example using AES in CTR mode
- Application: UUIDs

Stream Ciphers

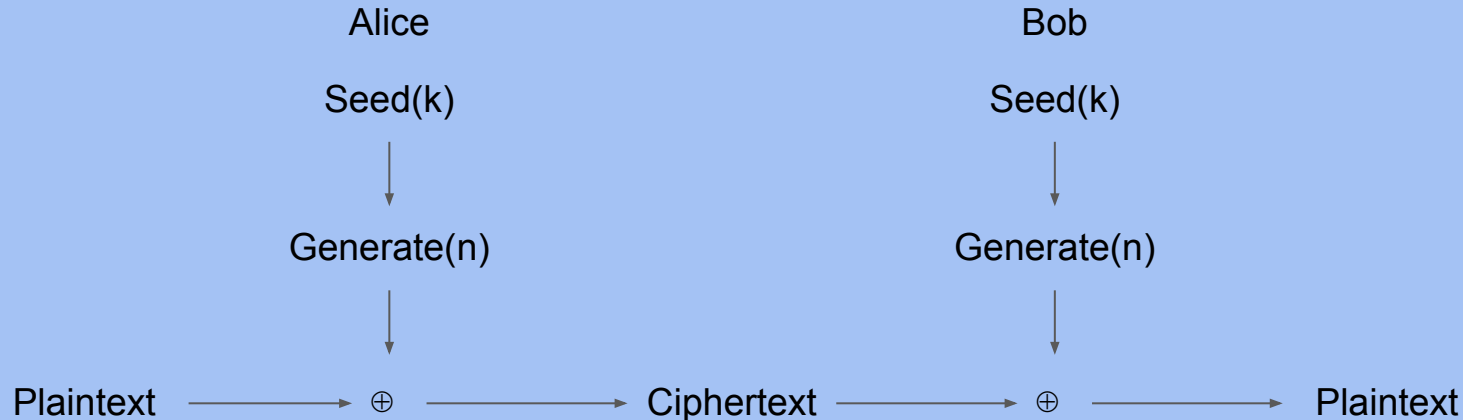
Computer Science 161

- Another way to construct symmetric key encryption schemes
- Idea
 - A secure PRNG produces output that looks indistinguishable from random
 - An attacker who can't see the internal PRNG state can't learn any output
 - What if we used PRNG output as the key to a one-time pad?
- **Stream cipher:** A symmetric encryption algorithm that uses pseudorandom bits as the key to a one-time pad

Stream Ciphers


Computer Science 161

- Protocol: Alice and Bob both seed a secure PRNG with their symmetric secret key, and then use the output as the key for a one-time pad



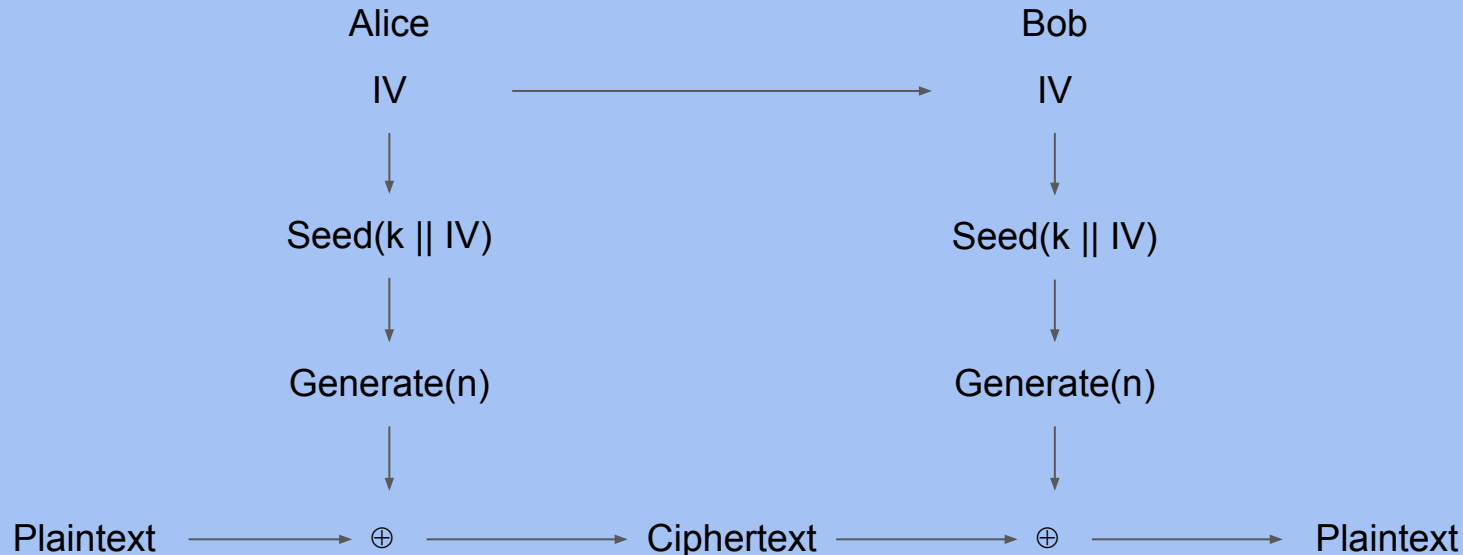
Stream Ciphers: Encrypting Multiple Messages

- Recall: One-time pads are insecure when the key is reused. How do we encrypt multiple messages without key reuse?



Stream Ciphers: Encrypting Multiple Messages

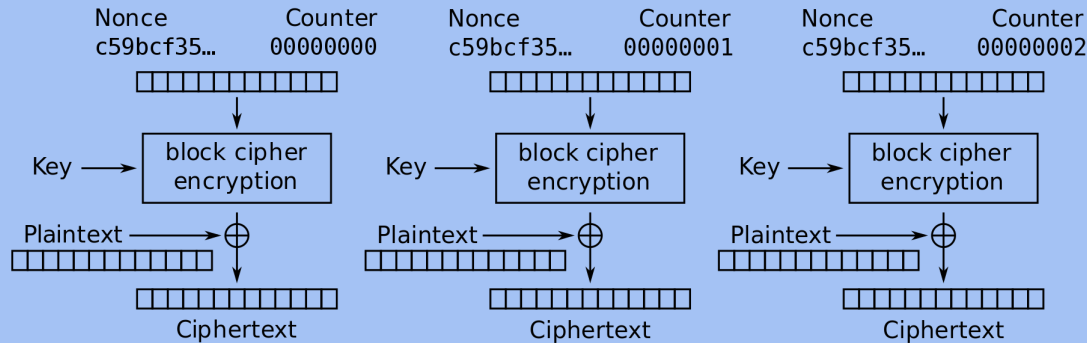
- Solution: For each message, seed the PRNG with the key and a random IV, concatenated. Send the IV with the ciphertext



Stream Ciphers: AES-CTR

Computer Science 161

- If you squint carefully, AES-CTR is a type of stream cipher
- Output of the block ciphers is pseudorandom and used as a one-time pad



Counter (CTR) mode encryption

Stream Ciphers: Security

- Stream ciphers are IND-CPA secure, assuming the pseudorandom output is secure
- In some stream ciphers, security is compromised if too much plaintext is encrypted
 - Example: In AES-CTR, if you encrypt so many blocks that the counter wraps around, you'll start reusing keys
 - In practice, if the key is n bits long, usually stop after $2^{n/2}$ bits of output
 - Example: In AES-CTR with 128-bit counters, stop after 2^{64} blocks of output

Stream Ciphers: Encryption Efficiency

Computer Science 161

- Stream ciphers can continually process new elements as they arrive
 - Only need to maintain internal state of the PRNG
 - Keep generating more PRNG output as more input arrives
- Compare to block ciphers: Need modes of operations to handle longer messages, and modes like AES-CBC need padding to function, so doesn't function well on streams

Stream Ciphers: Decryption Efficiency

- Suppose you received a 1 GB ciphertext (encryption of a 1 GB message) and you only wanted to decrypt the last 128 bytes
- Benefit of some stream ciphers: You can decrypt one part of the ciphertext without decrypting the entire ciphertext
 - Example: In AES-CTR, to decrypt only block i , compute $E_K(\text{nonce} || i)$ and XOR with the i th block of ciphertext
 - Example: ChaCha20 (another stream cipher) lets you decrypt arbitrary parts of ciphertext
 - What about HMAC-DRBG? You have to generate all the PRNG output up until the block you want to decrypt

Diffie-Hellman Key Exchange

Textbook Chapter 10

Cryptography Roadmap

Computer Science 161

| | Symmetric-key | Asymmetric-key |
|---------------------------|--|---|
| Confidentiality | <ul style="list-style-type: none">● One-time pads● Block ciphers with chaining modes (e.g. AES-CBC) | <ul style="list-style-type: none">● RSA encryption● ElGamal encryption |
| Integrity, Authentication | <ul style="list-style-type: none">● MACs (e.g. HMAC) | <ul style="list-style-type: none">● Digital signatures (e.g. RSA signatures) |

- Hash functions
- Pseudorandom number generators
- Public key exchange (e.g. Diffie-Hellman)

- Key management (certificates)
- Password management

Discrete Log Problem and Diffie-Hellman Problem

- Assume everyone knows a large prime p (e.g. 2048 bits long) and a generator g
 - Don't worry about what a generator is
- **Discrete logarithm problem (discrete log problem):** Given $g, p, g^a \bmod p$ for random a , it is computationally hard to find a
- **Diffie-Hellman assumption:** Given $g, p, g^a \bmod p$, and $g^b \bmod p$ for random a, b , no polynomial time attacker can distinguish between a random value R and $g^{ab} \bmod p$.
 - Intuition: The best known algorithm is to first calculate a and then compute $(g^b)^a \bmod p$, but this requires solving the discrete log problem, which is hard!
 - Note: Multiplying the values doesn't work, since you get $g^{a+b} \bmod p \neq g^{ab} \bmod p$

Discrete Log Problem and Diffie-Hellman Problem

Computer Science 161

For a random a, b, R :

$$g, p, \quad g^a \bmod p, \quad g^b \bmod p, \quad g^{ab} \bmod p$$

\sim  Indistinguishable from the perspective of a polynomial time attacker

$$g, p, \quad g^a \bmod p, \quad g^b \bmod p, \quad R$$

Diffie-Hellman Key Exchange

Computer Science 161

